


The London Interdisciplinary School
Master of Arts and Sciences in Interdisciplinary Problems and Methods
Capstone Project

What is AI slop?

Gustavo Araujo Costa



London, United Kingdom
July 15, 2025

Abstract

This project explores the emerging phenomenon of AI slop—low-effort, low-value AI-generated content that clutters digital information environments. While not necessarily false or malicious, slop is often meaningless, inauthentic, and aesthetically unappealing. Despite its growing visibility, the concept remains largely undefined in scholarly literature. To address this gap, I conducted qualitative (thematic analysis) and quantitative (topic modelling) analyses of sources referenced in the Wikipedia page on AI slop. I also utilized a large language model-based method, which offered comparable performance while being considerably more efficient. Thematic analysis reveals four core dimensions of the discourse around slop: efficiency and productivity, danger and discernment, epistemic crisis, and engagement economy. Slop is typically described as inaccurate, inauthentic, mediocre, surreal, distorted, generic, pointless, or kitsch—and is often metaphorically framed as flooding, contamination, infestation, or rubbish. Four main types of AI slop were identified: social media slop, corporate slop, submission slop, and slop news. This study also introduces the SlopNews dataset, used to analyse the statistical features of slop content. Compared to non-slop news, slop news is less complex, less varied, and tends to express more positive sentiment. These findings offer a foundation for understanding the dynamics and risks of low-quality AI content. They may inform educators, editors, researchers, and platform moderators in managing slop and safeguarding trust in digital communication.

Contents

1	Introduction	4
2	Critical cartography	9
2.1	Computer science	10
2.2	Philosophy, psychology and linguistics	11
2.3	Media studies	13
3	Thematic analysis	16
3.1	Methodology	16
3.1.1	Theoretical considerations	16
3.1.2	Data	16
3.2	Results	17
3.2.1	Efficiency and productivity	18
3.2.2	Danger and discernment	18
3.2.3	Epistemic crisis	19
3.2.4	Engagement economy	19
3.3	Qualitative analysis	21
3.3.1	Qualities	21
3.3.2	Metaphors	22
3.3.3	Types	22
4	Topic Modelling	25
4.1	Gensim	25
4.2	BERTopic	25
4.2.1	Interpretation approach	26
4.2.2	Filtering approach	26
4.3	LLM-based thematic modelling	28
4.4	Discussion	29
4.4.1	Note on LLM-based methods	30
5	Linguistic and statistical features	31
5.1	SlopNews: a dataset of AI slop	31
5.2	Analysis	32
5.3	Results	33
5.4	Limitations and future work	37
6	Conclusion	39
A	Appendix: Articles about slop	49
B	Appendix: Slop-related incidents	51
C	Appendix: BERTopic dendrogram	54
D	Appendix: LLM-based thematic modelling	56

1 Introduction

Every year, the Oxford University Press selects a Word of the Year that reflects their perception of the cultural landscape. The word chosen for 2024 was *brain rot*.¹ The shortlist also included the words *demure*, *dynamic pricing*, *lore*, *romantasy*, and *slop*.

They define *slop* as:

(*n.*) Art, writing, or other content generated using artificial intelligence, shared and distributed online in an indiscriminate or intrusive way, and characterized as being of low quality, inauthentic, or inaccurate.

Other Oxford Words of the Year – *podcast*, *selfie*, *carbon-neutral*, *post truth*, *rizz* – demonstrate their concern with online culture. *Slop* fits well into that tradition, as well as joining terms such as *fake news* and *enshittification* in the attempt to name a sense of increasingly pervasive “epistemic crisis” (Roberts, 2019) or “information disorder” (Wardle & Derakhshan, 2017).

Despite its growing cultural visibility, AI *slop*² lacks a formal academic definition. This study aims to fill that gap by exploring the meanings and practices associated with AI *slop*. In addition, it analyses statistical patterns in *slop* to identify its most salient features.

The problem of *slop*

As a general phenomenon, *slop* is not a new: humans have long generated “authentic” meaningless content for profit. Earlier forms include practices such as content farms, paper mills, clickbait, and search engine optimization (SEO) manipulation. However, recent developments in artificial intelligence (AI) technology have dramatically lowered the cost and effort required to create *slop*.

Consequences of this shift are already being felt in systems that rely on good-faith submissions, such as: open source coding projects (Purdy, 2025), job applications (Hoover, 2025), book publishing (Clarke, 2023), library catalogues (Maiberg, 2025), and academic papers (Tran, 2023). Across these domains, organizations are reporting a deluge low-quality synthetic content that they must manage (see Figure 1).

Slop also impacts systems which incentivise quantity over quality—such as engagement-driven social media and SEO. With increasing pressure to automate and scale content creation, *slop* has become a structural by-product of platform incentives and algorithmic visibility. Systems where the primary function is signalling effort – such as cover letters, recommendation letters and performance reviews – are expected to face similar disruption (Mollick, 2025).

The most conspicuous kind of AI *slop* is *social media slop*—the most notorious example being “Shrimp Jesus” (see Figure 3). This is surreal AI-generated

¹<https://corp.oup.com/word-of-the-year/>

²For the purposes of this work, I use “*slop*” and “AI *slop*” interchangeably.

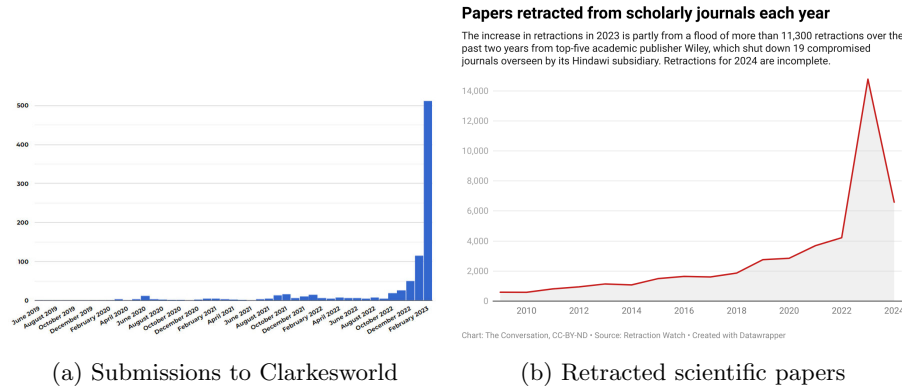


Figure 1: Sharp increases in low-quality or problematic submissions following the release of ChatGPT in 2022. Adapted from Clarke (2023) and Labbé et al. (2025a).

content spread on Facebook for driving engagement (Placido, 2024). Other examples include nonsensical phrases such as “vegetative electron microscopy” in fraudulent scientific papers (Joelving, 2025), fake Amazon product listings, (Tangermann, 2024), fake summary and biography e-books (Read, 2024), misleading AI-generated advice such as eating rocks (McMahon & Kleinman, 2024), and incorrect scientific illustrations such the infamous giant rat penis (Bailey, 2024).

Slop has been related to the “epistemic crisis” (Roberts, 2019), which is defined by misinformation, lack of trust in institutions, mainstreaming of conspiracy theories and erosion of shared reality. By flooding information environments with inauthentic, incoherent, or manipulative content, slop worsens epistemic instability, undermines the utility of digital systems, and places new burdens on already overstretched institutions.

Slop fits into the category of “information disorder” (Wardle & Derakhshan, 2017) alongside misinformation and disinformation. As Lazer et al. (2018) argue in the context of fake news, addressing such issues requires a multidisciplinary approach. This study contributes to that effort through the interdisciplinary integration of qualitative and quantitative methods with the objective of better understanding AI slop.

Methods and contributions

This study draws on two different data sources. First, the content of the references made in the Wikipedia entry for AI slop was analysed. For this purpose, three complementary methods were employed: thematic analysis (manual identification of themes), topic modelling (computational theme extraction), and LLM thematic modelling (LLM-based theme extraction). Apart from themes, the qualities, metaphors and types of slop are also analysed (see Figure 5).




Figure 2: Inaccurate scientific illustrations with nonsensical text. Adapted from Labbé et al. (2025a) and Bik (2024).



Figure 3: Shrimp Jesus.




Figure 4: Video thumbnails for get-rich-quick schemes using AI. Adapted from Vincent (2023).

Second, I introduce *SlopNews*, a dataset of slop news. It was constructed by scraping websites flagged by information reliability organizations as producing slop. These were combined with non-slop sources from publicly-available news datasets. Linguistic metrics were obtained and statistically analysed, finding that slop news exhibits less linguistic variation, less complexity and higher positive sentiment in comparison with non-slop news.

This work addresses pressing needs in academic integrity, content moderation, and information reliability. It provides theoretical and empirical foundations for those developing tools to curate the digital information landscape. While the subjective and evolving nature of AI slop presents some challenges, a rigorous, interdisciplinary approach is essential to mitigate its harmful effects and preserve epistemic integrity.




Figure 5: Dimensions (themes, qualities, metaphors and types) of AI slop.

2 Critical cartography

To better understand *slop*, I analyse related concepts and work in several disciplines—computer science, philosophy, psychology and media studies, as well as the origins of the term itself. It is a “critical” cartography in the sense that it uses attrition with its source material to create insight (instead of merely reporting it). It does not try to maintain a neutral and objective view of the content but leverages its affective component in favour of the analysis.

Although the word *slop* was used to refer to low-quality writing as far back as the 19th century (“Oxford Word of the Year 2024”, [n.d.](#)), coinage of its current AI-related usage has been attributed to Willison ([2024](#)). According to Aleksic ([2024](#)), two key events are related to its popularisation: the creation of its Wikipedia page and the publication of an article in the New York Times (Hoffman, [2024](#)). From its origin in AI and tech internet circles, the term’s usage has seen a remarkable 334% increase in usage in 2024 (“Oxford Word of the Year 2024”, [n.d.](#)).

While no academic papers on AI *slop* have been published, there are two preprints: an unrelated technical advance in AI writing quality (Chakrabarty et al., [2025](#)) and a discussion of AI-generated propaganda, which they call *slopaganda* (Klincewicz et al., [2025](#)). While the scientific literature is sparse, there are numerous op-eds, blog posts and similar media which extensively discuss the subject.

Slop is similar to other forms of problematic information:

- **Disinformation:** false information spread with malicious intent (Wardle & Derakhshan, [2017](#)).
- **Misinformation:** false information spread inadvertently (Wardle & Derakhshan, [2017](#)).
- **Hallucination:** AI-generated content which seems plausible but is factually incorrect (Huang et al., [2025](#)).
- **Rumour:** information which is unverified; a debunked rumour can be classified as misinformation or disinformation depending on the intent of the source (Zubiaga et al., [2019](#)).
- **Fake news:** news which are false and public; unlike rumours, they always refer to facts which can be verified (Zubiaga et al., [2019](#)).
- **Spam:** Intrusive messages typically associated with scams.
- **Bullshit:** communication characterized by indifference to truth (Frankfurt, [2009](#)).

Unlike misinformation and hallucination, *slop* is not necessarily false. Unlike spam and disinformation, it is not necessarily malicious. Its closest parallels are spam (with its emphasis on quantity) and bullshit (with its emphasis on plausibility). *Slop* has been likened to a metaphorical “distributed denial-of-service” (DDoS) attack, where an excess of useless requests prevents a system from processing legitimate ones (Purdy, [2025](#)).

2.1 Computer science

In natural language processing (NLP), the classification of problematic information typically involves linguistic pattern analysis, statistical methods, and fact verification (Aich et al., 2022; Tang et al., 2023). However, the highest performance is offered by fine-tuned deep learning models, which don’t require manual feature creation at all (Chakrabarty et al., 2025; Chen & Shu, 2024; Doughman et al., 2024; King et al., 2024). The developers of Pangram offer a detailed technical report on their deep learning methodology (Emi & Spero, 2024). Full Fact, a fact-checking organization, also uses deep learning to prioritize rumours which are likely to spread.³

Widely-used commercial AI detectors such as GPTZero⁴ and Originality.ai⁵ employ various features to distinguish human from machine-generated text. These metrics include perplexity, burstiness, readability, and patterns in punctuation and word frequency. Tools like the Problematic Paper Screener (Labbé et al., 2025b) and Wiley’s papermill detector (Flanagan et al., 2024) identify “fingerprints”, which are telltale phrases and artifacts left by software).

It remains unclear whether detecting slop is significantly different from detecting AI-generated content. Given that slop is defined by its lack of effort, AI detectors should be able to filter it effectively. However, this would penalize high-quality AI-generated content, which is, presumably, not a desirable outcome. Additionally, AI-detection tools are notoriously unreliable (Forlini, 2023; Sadasivan et al., 2023; Weber-Wulff et al., 2023) and lack interpretability, which is essential for algorithmic accountability (Kapoor et al., 2024; Saphra & Wiegrefe, 2024; Templeton et al., 2024). Although interpretability methods exist, these often yield suboptimal or misleading results (Gonen & Goldberg, 2019; Heap et al., 2025). This suggests that a simple, feature-based slop classifier could be a promising direction for future research.

Another method devised for identifying misinformation involves simply prompting LLMs for analysis (Chen & Shu, 2024). This shows remarkable results compared to other approaches. It also is more interpretable, since the LLM output comes with an explanation.⁶

NLP classification tasks can be broadly divided into those with objectively verifiable answers (e.g., fake news, AI-generated text) and those where ground truth is ambiguous or subjective (e.g., rumours, hate speech). AI slop falls into the latter category, as its fundamental characteristics are not grounded in objective criteria. However, even in cases where factual verification is possible, fact-checking is labour-intensive and time-consuming. As a result, fake news is often addressed by labelling the credibility of the source rather than analysing the content itself (Peters, 2022). Another approach involves using machine-learning algorithms to quickly determine which rumours are likely to spread, instead of using them for the actual classification (Zubiaga et al., 2019).

³<https://fullfact.org/ai/about>

⁴<https://gptzero.me/>

⁵<https://originality.ai/>

⁶Although the explanations are, strictly speaking, hallucinated.

AI risk taxonomies

There are many taxonomies of AI risk currently being developed. AI slop does not fit neatly into established risk categories, since it neither causes physical harm nor is obviously toxic. At an individual level, its effect is only inconvenience. Taken as a collective phenomenon, however, it poses psychological, economic and political risks by undermining trust in the information environment.

In the MIT AI Risk Taxonomy (Slattery et al., 2025), constructed from other 65 other taxonomies, slop is most closely aligned with subdomain 3.2: “pollution of information ecosystem and loss of consensus reality”. Although this category is primarily concerned with issues like personalization and filter bubbles (also known as echo chambers), it remains the most fitting descriptor for the risks associated with the proliferation of slop.

Another relevant taxonomy was developed by the Center for Security and Emerging Technology (Hoffmann & Frase, 2023). In this framework, we can categorize slop as an intangible issue of detrimental content and/or democratic norms. Here, *intangible* refers to phenomena that are difficult to observe or lack a direct material manifestation, while *issue* denotes an ongoing concern or a possible occurrence rather than a specific event. Slop falls under detrimental content through its association with misinformation, and it relates to democratic norms via its contribution to information disorder.

2.2 Philosophy, psychology and linguistics

Philosophical discourse has extensively explored ideas of the “loss of humanity” brought about by technological advancement: from Plato’s concerns about writing and memory (Derrida, 2020), to Heidegger’s critique of modern technology (Graña, 2016), and now contemporary philosophers of different traditions (Han, 2015; Krenak, 2020).⁷

However, the defining feature of slop is not its artificial origin, but its poor quality. This is similar to hallucination, where the issue is not artificiality, but inaccuracy. Artificial content which is truthful and high-quality is rarely considered slop, as it is not recognized as artificial or inauthentic at all. This also poses a challenge for the collection of high-quality AI-generated content, as it is indistinguishable from high-quality human-made content.

This distinction can be observed in the Google Search’s guidance for developers, where they state their interest on “the quality of content, rather than how content is produced” (Sullivan & Nelson, 2023)⁸. This highlights the importance of three different interlocked parts: the origin of the content, the content itself, and the interpretation of the reader.

⁷Although there are alternative accounts, such as in science and technology studies (Haraway, 2015)

⁸In this document, they warn that using automation to manipulate search results violates their spam policies.

With the advent of AI as a general-purpose technology, novel interdisciplinary work is emerging. One example is the exploration of LLMs in the context of the philosophical concept of bullshit (Gorrieri, 2024; Hicks et al., 2024; Trevisan et al., 2024). Bullshit is defined as discourse indifferent to its truth-value; instead of prioritizing accuracy, it is created for some other objective, such as plausibility or persuasion (Cohen, 2002; Easwaran, 2023; Frankfurt, 2009). This is distinct from lying, which involves intentionally stating a known falsehood, and thus having an interest in its truth-value.

Instead of “hallucinating”, AI systems can be better understood to be bullshitting – generating plausible-sounding text without regard for truth (Hicks et al., 2024). In this sense, generating slop is also bullshitting: communicating something not for its informational value but for its effect (e.g., engagement on social media).

The traditional definition of bullshit centres on the intentions of the speaker. There is an alternative definition, which defines it exclusively through its contents. According to this view, bullshit is defined by being ‘unclarifiable’ (Cohen, 2002). This is the definition most commonly used in psychology, particularly in studies of susceptibility to “pseudoprofound bullshit” (Pennycook et al., 2015) and fake news. This approach deliberately sidesteps the philosophical emphasis on the speaker and instead focuses on a property of the content itself. In this sense, it parallels Google’s policies which prioritize the quality of the content over its authorship.

Along with the relationship between author and content, we can also consider the one between content and consumer. In the context of the subjective judgment of “quality”, it is important to recognise that beauty is in the eye of the beholder—what one person perceives as inauthentic, another might find meaningful (Dalton, 2016). This emphasizes the role that normativity and context play in subjective labels such as bullshit and slop. This can be observed in the motivations for the creation of the different definitions of bullshit: the original intent-based definition was created to describe rhetoric in politics and advertising; the content-based view was developed to expose obscurity in French Marxism; and psychological approaches were created to criticise new-age spirituality. In their attempts to objectively define the concept (by moving away from the origin and the interpretation towards analysing the content itself), the researchers ultimately only make their biases more prominent. Informed by this dynamic, for my analysis, I select sources which were previously determined to be slop and whose intentions seem to be uncontroversially questionable.⁹

The psychological study of epistemically suspect beliefs – such as bullshit and fake news – draws heavily on the heuristics-and-biases tradition from behavioural economics (Tversky & Kahneman, 1974). This cognitive perspective is central to the analysis of slopaganda in Klincewicz et al. (2025), where susceptibility to slop is related to phenomena such as the illusory truth effect (where repeated exposure increases perceived truthfulness), confirmation bias

⁹This raises the possibility of studying bullshit through statistical properties, though this would require controlling for stylistic differences.

(the tendency to favour information that confirms existing beliefs), and information overload (where emotionally charged content captures limited attention resources).

This body of literature, however, tends to frame bias as the exploitation of vulnerabilities in human rationality (for example, see Muzumdar et al., 2025; Pennycook & Rand, 2019). In this way, it neglects the performative and emotional dimensions of communication—treating misinformation as a failure in the transmission of information rather than a socially and affectively embedded practice (Wardle & Derakhshan, 2017). Similar critiques have already been raised in economics (Infante et al., 2016) and sociology (Barker, 2011), where reduction of behaviour to rational calculation and information transfer is seen as inadequate.

The interpretation of text involves more than just communicating information. According to relevance theory (Sperber & Wilson, 1987), communication depends on attention and context: information is considered relevant only when it connects to certain implicit assumptions. This process requires considerable inference and intention. A “rational communicator” is presumed to cooperate in making such relevance accessible.

Slop violates this expectation by being irrelevant. While obscure texts may still signal some underlying authority or value – what Sperber (2010) calls the “Guru Effect” – slop offers neither. Rather than inviting interpretation, slop invites disengagement. It does not gesture toward deeper meaning or hidden insight; it is low-effort noise.

2.3 Media studies

Rather than treating slop as a purely technological or psychological artifact, I also consider the structural and socioeconomic conditions that incentivize slop production: algorithmic visibility, low-cost scaling, and platform monetization.

There is substantial public anxiety surrounding misinformation (Cellan-Jones, 2017; World Economic Forum, 2024). In the United States, adults perceived made-up news as a more pressing issue than violent crime, climate change, racism, illegal immigration, terrorism, or sexism. A survey found that 64% of US citizens believe “fabricated news stories cause a great deal of confusion about the basic facts of current issues and events” while 51% “cite the public’s ability to distinguish between facts and opinions as a very big problem” (Barthel et al., 2016). The rise of AI-generated content – which has been shown to be highly persuasive in certain situations (Costello et al., 2025; Koebler, 2025) – risks making this situation worse.

AI slop operates similar to disinformation (Wardle & Derakhshan, 2017) in that both consist of repetitive and low-value content. This is remarkably similar to the political strategy articulated by Steve Bannon: “flood the zone with shit” (Illing, 2020). The goal is not to convince, but to exhaust and confuse. While slop does not have these as its main purposes, they happen as consequences of the prevalence of slop in an informational environment.

Ironically, information about misinformation is often misleading. Evidence suggests the problem posed by fake news is overstated. The values cited above reflect the public *perception* of the issue, not its empirical prevalence. The average American’s exposure to fake news has been found to be low, at around 0.15% (Allen et al., 2020). A study by the UK Centre for Emerging Technology and Security found (CETaS) found no evidence that AI disinformation or deepfakes impacted UK, French or European elections results in 2024 (Stockwell, 2024). In most cases, social media actually *increases* exposure to alternative viewpoints, while only about 8% of the population exist in ideological “echo chambers” (Arguedas et al., 2022). This reveals that the core danger is not necessarily an information environment “poisoned” with falsehoods, but the pervasive fear and corrosion of public trust that the perception of such an environment creates.

Organizations tackling this issue, however, are still subject to the pressures that generate misinformation. Vendors of technical solutions and academics alike lean into the perceived risk of “information disorder”. By aligning with grand narratives about AI’s capacity to either save or destroy humanity, they mystify the technology, which helps maintain its allure and its perceived threat (Carpenter, 2025; Tully et al., 2025) – thus maintaining their own relevance in an information environment characterized by scarcity of attention and abundance of noise.

Therefore, the goal should not be to simply “neutralize the poison” with technical fixes, but rather to rebuild the perception of trust in our information environment. Misinformation is a complex social problem that resists simple technological solutions (Singh, 2024). While tools like detection algorithms can help, they must be part of a broader, interdisciplinary strategy.

Enshittification

Market incentives – from click-driven ad revenue to “publish or perish” academia – fuel the production of slop by motivating agents to create a high volume of low value content (Knibbs, 2024a; Labbé et al., 2025b). In this way, this is not an issue of psychological susceptibility, but of the consequences of the systems with which people are compelled to engage.

In this work, I use the term “epistemic crisis” to refer to these generalized concerns about misinformation (e.g. vaccine hesitancy during the pandemic), disinformation (e.g. Russian propaganda about Ukraine), conspiracy theories (e.g. QAnon during the US election), pseudoscience (e.g. hydroxicloroquine as a treatment for Covid-19), and related phenomena. This dynamic is particularly evident on platforms like Facebook, which has a troubled history of managing political influence, as seen in cases like Myanmar and Cambridge Analytica (Wynn-Williams, 2025).

Digital platforms play a large role in the dynamics of the epistemic crisis. One interesting concept to understand their functioning is *enshittification* (Doctorow, 2023). It refers to the process by which platforms (e.g., Amazon, TikTok) compromise the quality of their service. First, they attract users with high-quality services while operating at a loss. Once users are locked in, the

platform shifts its focus to extracting maximum value for shareholders, systematically degrading the user experience in the process.

Similar to slop and bullshit, the term is sometimes used as a “general term of abuse” (Frankfurt, 2009), that is, as a strictly derogatory term. Enshittification is often used to refer to a generalized degradation of systems. In contexts where plausibility and engagement are prioritized over accuracy and “connection,” the result is what we might call “slopification”—systematic incentives toward the creation of low-effort and low-quality crap instead of things which are actually worthwhile. Through the execution of this study, I became convinced that we have moved on from the post-modern condition (Lyotard, 1984) to a new cultural moment: “we live in the slop era” (LowercaseJai, 2025).

3 Thematic analysis

3.1 Methodology

3.1.1 Theoretical considerations

To examine how the term “AI slop” is used, I employ thematic analysis—a widely adopted qualitative method for identifying patterns of meaning across a dataset through manual coding and interpretation (Braun & Clarke, 2006; Squires, 2023). This approach was selected for its flexibility and depth, which are well suited to the exploratory study of an emerging and loosely defined concept. In practice, thematic analysis involves coding segments of text with themes that are inductively derived from the data themselves.

Following Braun and Clarke’s framework, I explicitly outline the project’s epistemological assumptions. This project attempts to be: inductive (themes arise from the data instead of an existing theory), semantic (focus on explicit data instead of latent meaning), and constructionist (focus on sociotechnical and political aspects of the themes instead of particular experiences). Theme prevalence is assessed in terms of how many distinct items reflect a theme, rather than how often the theme appears overall.¹⁰ In qualitative research, subjectivity is not treated as a bias to be eliminated but as an integral part of the interpretive process. As such, I acknowledge my position as a researcher who regularly uses AI tools, but understands them as “normal technology” (Narayanan & Kapoor, 2025)—neither utopian nor dystopian. Viewing all technologies as embedded in social systems, my goal is to understand how slop is used to fulfil a sociotechnical need.

3.1.2 Data

The dataset consists of 34 articles cited by the Wikipedia page on AI slop;¹¹ these were selected after discarding tangential academic preprints and repeated articles on the same topics. Content was collected using Trafilatura¹² and coded using Taguette.¹³ While I refrain from publishing the content of the dataset due to copyright concerns, a complete list of sources is included in Appendix A.

This dataset has several limitations:

- It represents a temporal snapshot and does not capture recent developments in the fast-moving phenomenon of AI slop;
- It reflects the editorial decisions of the Wikipedia contributors, which is evident in the prominence of topics such as advertisement and video games.
- While several articles criticise the incentives of the engagement economy, they are subject to the same attention-seeking dynamics. Newspapers are

¹⁰Prevalence is not reported in the body of the text due to practical constraints.

¹¹https://en.wikipedia.org/wiki/AI_slop

¹²<https://github.com/adbar/trafilatura>

¹³<https://www.taguette.org/>

incentivised to use emotional language and dramatize their narratives.¹⁴

As it is common for slop to be used as a catch-all term for dissatisfaction with AI-produced content, I deliberately did not distinguish between mentions of slop and more general commentary on AI.

3.2 Results

Table 1: Themes of AI slop.

Theme	Description
Efficiency and productivity	Slop is generated due to its low cost, speed, ease of creation, and scalability. This applies to corporations aiming to cut costs and individuals seeking side income. These indicate a significant incoming shift in several industries, which is accompanied by concerns about the automation of creative labour.
Danger and discernment	The excess of slop overwhelms information systems and creates the need for its filtering and moderation. This is especially challenging since slop does not strictly violate platform policies. This is threatening both to institutions and to individual users, which might unknowingly be exposed to false information.
Epistemic crisis	Slop is related to moral anxieties about political disinformation, the “post-truth” environment, and the wider breakdown of shared reality and trust in institutions.
Engagement economy	Platform incentives (such as engagement, virality, and ad revenue), which are ultimately about profit, are the primary drivers of slop. Individuals in low-income regions are especially incentivised to engage in its creation.

As shown in Table 1, four major themes were identified in the dataset: *efficiency and productivity*, *danger and discernment*, *epistemic crisis* and *engagement economy*. These are highly interrelated: AI’s speed and scale interacts with platform’s quantity-over-quality incentives, creating moderation challenges

¹⁴As so am I. This project made me very aware of the ways I can choose to lean into metaphorical, emotional or apocalyptic language to create a more compelling narrative — thus contributing to slopification.

and polluting the information environment. Some themes had surprisingly low prevalence, such as: general AI risks (model collapse, guardrails, copyright, environmental cost, existential risk, plagiarism), positive slop (e.g. as art or humour), and other hyped technologies (e.g., cryptocurrency, NFTs).

Each theme is described below. This section references the articles’ dataset ID (e.g., A5 for article 5) as shown in Appendix A.

3.2.1 Efficiency and productivity

It’s easy to learn, almost zero cost, and can be done any time at home with just a phone. (A12)

Ryanair and Samsung are jumping on the trend and seeing six-figure engagement spikes — all while barely lifting a production finger. (A21)

This theme is related to the reasons why people choose to produce slop. Sub-themes include: *cheap, easy, fast, scalable, consumable, productive*, and *transformation*. These qualities appeal to both corporations seeking cost-cutting and individuals (especially in low-income regions) seeking side income.

Generative AI increases productivity by expediting content creation —usually in visual media. One “slopper” comments that his method consists of “producing far more content than people who don’t use AI” (A12), while a corporate spokesperson emphasizes that more than cost-cutting, it is about the speed of production, saying that AI makes it “five times” (A7) faster.

This logic of scale rewards quantity over quality, producing the low-effort aesthetic associated with slop: “amateurish, really corny, really bad art” (A9). For corporations, the appeal lies in productivity gains: “developers can now make stadiums in six weeks rather than six months” (A29). But these shifts are contentious, often seen as threats to creative labour and artistic integrity: “here’s hoping Activision hires some actual humans” (A20). Usage of AI tools is usually discovered by encountering distorted media: “loading screens depicted hands with extra digits” (A30).

These transformations are understood to be significant, as “it’s clear that business as usual won’t be sustainable” (A5) and that it is just the start of the process—“so it begins” (A22).

3.2.2 Danger and discernment

But the slop tide threatens some of the key functions of the web, clogging search results with nonsense, overwhelming small institutions [...], and generally polluting the already fragile information ecosystem of the internet. (A5)

This theme is related to the consequences of the abundance of slop and the strategies used to deal with it. Subthemes include *danger, spam, scam, wading*

through, moderation, filtering, detection, and enforcement. The excess of AI-generated content (Theme 1) creates the need for a strategy to deal with it. Libraries, search engines, and platforms struggle to filter growing volumes of synthetic content: “overwhelmed and underfunded local libraries” now house “unchecked, unedited AI-generated artifacts” (A5).

Unlike false information and malicious campaigns, slop is not overtly harmful: “[it’s] helpful to think of these photos and illustrations not as nefarious deepfakes or even hyper-persuasive propaganda, but as digital chum” (A14). Slop doesn’t “technically violate any rules” (A3), which complicates its enforcement. Platforms might tolerate it because “the slop makes everyone money” (A13). While not directly dangerous, “slop is often scam-adjacent” (A3) and is referred to as “threatening” (A7). For example, there are slop mushroom picking books which “contained dangerous advice for anyone hoping to discern a lethal fungus from an edible one” (A2), and chatbots can questionable advice on “potentially delicate subjects” such as chemotherapy and Ozempic (A6).

3.2.3 Epistemic crisis

AI slop more broadly is poisoning the internet and is destroying the last shred of any sense of reality or reliability of social media platforms and search engines. (A15)

This theme is related to slop’s large scale social impact, such as its relationship to politics and misinformation. Subthemes related to this include *politics, conspiracy, reality, trust, and trump*. It relates to the epistemic crisis—a cultural climate of lack of trust in institutions, political polarization, worry about fake news and erosion of a shared reality. Trump, and the aesthetics of the alt-right, are repeatedly referenced as emblematic: “Schoolcraft’s Facebook group is full of examples of troop-saluting, anti-LGBTQ+, pro-Trump schlock” (A9).

This breakdown of shared meaning and trust escalates the concerns about discernment (Theme 2) into more abstract and existential territory. While there are material concerns (e.g., preserving the credibility of advertisement organizations; A7), the dominant framing shifts towards moral and metaphysical anxieties. Truth and trust are depicted through metaphors of pollution and decay. Slop is not treated as merely an issue of quality, but as a hallmark of the post-truth condition: “we live in an era where the truth essentially does not matter” (A15); “the far right’s world-building project, where feel is always greater than real” (A16).

3.2.4 Engagement economy

This reality-fracturing is the result of an information ecosystem that is dominated by platforms that offer financial and attentional incentives to lie and enrage (A16)

Just like spam, almost no one wants to view slop, but the economics of the internet lead to its creation anyway. (A2)

This theme tells the story of slop as a consequence of the incentive structures that organize the internet. Subthemes include: *platforms, corporations, engagement, enshittification, incentives, money, virality, developing countries*. Central to this narrative is the influence of algorithmic systems, which shape content production in ways that prioritize reach and profitability over quality. As a result, slop is frequently linked to individuals producing viral material as part of informal economies—described as “side hustles” (A12) or “trying to make a quick buck” (A14).

Articles frequently emphasize the role of “incentives”, suggesting that slop is not primarily the result of personal proclivities, but structural forces. Instead of providing meaningful experiences, sloppers are prioritizing metrics to maximize revenue. For this purpose, “the content of the content isn’t really as important as its presence—or, more accurately, its measurability” (A5). Facebook’s Performance Bonus Program – referred to by one article as a “slop subsidy” (5) – rewards creators based on the engagement of their posts, which ultimately results in the creation of social media slop: viral nonsense posts with attention-grabbing surreal imagery (e.g., Shrimp Jesus).

Although the term enshittification is not explicitly mentioned, its logic permeates much of the discussion, evident in phrases like “the great internet let-down” (A12), or:

And so it begins – the absolute dumbing down and worsening of everything from media to film to advertising to politics to the written word at the dead hands of AI. (A22)

While Theme 3 reflects metaphysical stakes (truth, trust, authenticity), Theme 4 reveals geopolitical ones (capital, platform governance, digital colonialism). Articles often refer to slop operations in developing countries — Kenya (A5), Vietnam (A11), Pakistan (A13), Malaysia (A15), etc. At the centre of this network are money and advertisement: money made by slop producers through platforms, by platforms through advertising, and by advertisers through increased productivity. As one article states, “AI slop is being generated by people all over the world for the express purpose of going viral and making money” (A15).

3.3 Qualitative analysis

This part of the study expands the analysis of AI slop beyond thematic analysis into a broader landscape of qualities, metaphors, and types. This expansion is motivated by the richness of the data (which did not fit neatly into methodological boundaries), the aim for thoroughness (given the study’s introductory nature), and the practical value of reporting on the multidimensional nature of slop.

3.3.1 Qualities

Table 2: Qualities of AI slop

Quality	Related words
Excessive	<i>everywhere, clutter, spam, inescapable, river, vast quantities, churn, glut, fatigue, high-volume, unmanageable, inescapable</i>
Inaccurate	<i>wrong, hallucinated, incorrect, dangerously inadequate, misleading, nonfactual, cartoonishly false, inaccurate, fake</i>
Inauthentic	<i>lifeless, voiceless, mawkish, inhuman, meaningless, authentic, legitimate, intent, purposeless, unnatural, mindlessly</i>
Mediocre	<i>low-quality, low-grade, shovelware, low-effort, too easy, cheap, shoddy, trivial, unskilled</i>
Surreal	<i>surreal, uncanny, bizarre, unsettling, outlandish, gobbledygook, nonsense, disturbing, dreamscapes</i>
Distorted	<i>distorted, warped, deformed</i>
Generic	<i>generic, filler, repetitive, unoriginal, unimaginative, banal, flat, boring, familiar, flat affect</i>
Pointless	<i>plausible, inconvenient, inoffensive, irrelevant, contextless, unneeded, useless</i>
Kitsch	<i>kitsch, shlock, high-resolution, low-fi, loud, messy, amateurish, corny, bad art, cringe, lowbrow, manipulative</i>

The outcome of the proliferation of slop is the drowning out of intention, meaning, and effort; words like *intentionality*, *depth*, *creativity* and *value* are often invoked as slop’s opposites.

Some terms didn’t consolidate into qualities but recur: *disrespectful*, *inconvenient*, *inoffensive*, *good*. Many qualities – like *kitsch*, *surreal*, *excessive* – tie back to “weird AI crap” (A5) on Facebook, while others – *inauthentic*, *inaccurate*, *generic* – are more relevant to other types of slop such as scientific papers, e-books and news articles.

3.3.2 Metaphors

Metaphors are fundamental to how humans communicate (Lakoff & Johnson, 2008). They not only help communicate experiences but also shape experience, understanding and behaviour (Chkhaidze et al., 2021). In this context, we identify central metaphors used to talk about slop.

Table 3: Metaphors

Metaphor	Relevant words
Flooding	<i>flood, clogging, deluge, wade through, stream, drowning, rising tide, swamp, torrent, outpouring, dilutes, floodgates, river, drool</i>
Contamination	<i>pollution, scrubbing, ecological footprint, chokes</i>
Infestation	<i>spreading, proliferation, propagation, grow wildly, mulch, weeding out</i>
Rubbish	<i>trash, rubbish, garbage, dross, crap, junk</i>

These metaphors show themes of flooding (water and overabundance), contamination (pollution and toxicity), infestation (unchecked spread and infiltration) and rubbish (worthlessness and clutter). In these we find the emotional aspects of slop: its dirtiness, excess, vacuity and threat. Among these, the most dominant idea is *flooding*.¹⁵

Language about slop is not only metaphorical, but displays a characteristic sound symbolism. Words related to slop share a kind of sludgy phonosemantics: mulch, glut, churn, chum, spam, etc. Other related concepts that share metaphorical aspects of excess, inauthenticity, and intrusiveness include “pink slime” (inauthentic media content), “slush pile” (unsolicited manuscripts), “vendor slurry” (library slop) and “grey goo” (self-replicating nanobots).

Slop also has a kinetic quality. Through the water metaphor, it is related to flowing water, or to clogging information flow. Through infestation, it spreads—undesirable, everywhere. Other minor metaphorical resonances identified include: inconvenient but not threatening (“clutter”), mechanical or inhuman (“churn out”, “noise”), vacuity (“drivel”), excess (“glut”) and undesirable food (“unappetizing”).

3.3.3 Types

While social media slop is easily recognisable, there are other facets of the same phenomenon. These include:

¹⁵Flooding has also been found in political contexts as a dehumanizing way to refer to immigrants (Mendelsohn & Budak, 2025)

- **Social media slop:** low-effort or manipulative engagement bait created by sloppers for profiting on social media platforms. This kind of slop is primarily identified by its *surreal* and *kitschy* appearance.
 - **Example:** “The stuff of Salvador Dali’s wildest dreams is no match for Facebook these days: Amputee kittens using crutches. Strawberries in the shape of lifelike frogs. Bosomy conjoined twins, structurally impossible sand sculptures, snakes swallowing fully-grown lions, airplanes with human hands. An underwater Jesus covered in shrimp.” (A9)¹⁶
- **Corporate slop:** low-quality AI-generated media used by corporations for cost-cutting purposes, typically in advertising. This kind of slop is primarily identified through its *distorted* quality, signalling *inauthenticity* and *lack of effort*.
 - **Example:** “A thumbnail image for 12 Angry Men on Amazon Freevee used AI to depict 19 men with smudged faces, none of whom appeared to bear any similarities to the characters in the film.” (A1)
- **Submission slop:** papers, books, applications and other media which are mass-submitted and overwhelm selection systems. This kind of slop is identified by its *excess* which overwhelms systems, its *inaccurate* information within media where accuracy is essential, and by being *generic* or *mediocre* in other contexts.
 - **Example:** “identical Viking ‘novels’ with seemingly AI-generated covers, all called Wrath of the Northmen: A Gripping Viking Tale of Revenge and Honor (that one has been published variously by authors named Sula Urbanr, Sula Urbanz, and Sula Urbanr).” (A5)
- **Slop news:** low-quality news articles created by professional sloppers for SEO manipulation instead of human readership. This kind of slop is identified by being *generic* and *pointless* in comparison with news that was written for human consumption.
 - **Example:** see Appendix E

This is not an exhaustive list. Other examples could include summarization hallucinations – like “Google suggesting that you could add nontoxic glue to make cheese stick to a pizza? That’s slop” (A8) – as well as politically motivated disinformation, such as “Kamala Harris did not give a speech at the Democratic National Convention to a sea of communists while standing in front of the hammer and sickle” (A14). However, these cases are more appropriately addressed under established categories (hallucination and disinformation, respectively). Similarly, the fake events prevalent in this dataset resulted from low-effort SEO hacking, and are not necessarily related to AI.

¹⁶For examples of social media slop images, see the [Ai Boomertrap\(TM\)](#) Facebook group and the [Insane Facebook AI Slop](#) subreddit.

Table 4: Types of slop.

Type	Purpose	Product	Media	Creator	Consumer	Primary qualities
Social media slop	Engagement	Social media content	Image	Sloppers	Users	Surreal, kitsch, pointless
Corporate slop	Productivity	Ads, films, games	Image, audio, video	Corporations	Users	Distorted, mediocre, inauthentic
Submission slop	Submission	Books, papers	Text, image	Sloppers, paper mills	Publishers, journals	Excessive, inaccurate, mediocre
Slop news	Ad revenue, SEO	News articles	Text, image	Sloppers	Users and bots	Pointless, generic, mediocre

Slop is often economically motivated, closely tied to advertising, cost-cutting, propaganda, and scams. Advertising, in particular, plays a central role. Prominent examples include corporations generating AI-produced ads and assets, as well as individual “sloppers” creating viral images for engagement.

Visually, slop is most recognizable in image-based content. This is likely due to the ease with which artifacts – distorted anatomy, surreal ideas – can be identified, particularly in social media and corporate contexts. The prevalence of slop in text is unknown, as it is harder to detect—something we investigate in the last section of this study.

4 Topic Modelling

Interdisciplinary research involves the integration of different methods as a tool for more insight into its object of study. For this purpose, we employ *topic modelling*—an NLP technique which, similar to thematic analysis, aims to identify themes (topics) within a dataset. The libraries utilized for this were **gensim** and **bertopic**. We also reincorporated the articles which were previously discarded due to redundancy.

4.1 Gensim

Gensim is a Python topic modelling library. We applied a standard topic modelling technique utilizing TF-IDF and Latent Dirichlet Allocation (LDA), setting hyperparameters to 10 passes, 5 topics, and auto alpha. Text underwent pre-processing, including tokenization, lemmatization, lowercasing, and removal of stop words, punctuation, and numbers. Results are reported in Table 5.

Table 5: Topic modelling results using TF-IDF and LDA.

Topic	Top 10 Words
0	enix, cola, coca, square, foamstars, voice, actor, okatani, album, vercel
1	parade, game, halloween, activision, slop, dublin, event, facebook, image, website
2	coca, cola, meta, advert, baby, maga, harris, bandicoot, click, facebook
3	claude, tweet, unwanted, term, ethic, llm, spam, aleksic, may, language
4	trump, cola, brand, coca, spotify, centric, italian, brainrot, mcdonald, hurricane

Topics 0 and 1 primarily discuss corporate slop in gaming and fake events, Topics 2 and 4 focus on corporate slop in advertising and politics, and Topic 3 focuses on AI slop. All topics lean heavily towards named entities. As these are not very insightful, we explore a more advanced model.

4.2 BERTopic

BERTopic streamlines a multi-step topic modelling pipeline. Documents are first segmented into sentences. From these, contextual embeddings are generated, reduced in dimensionality and clustered. For each cluster, a bag-of-words representation is created and topics are extracted. Finally, an LLM is used to name the topic. The number of topics can be adjusted through a hyperparameter of the clustering model. Following the documentation,¹⁷ we use:

¹⁷https://colab.research.google.com/drive/1BoQ_vakEVtojsd2x_U6-_x52OOuqrj2

- Embeddings: `text-embedding-3-large` (instead of `all-MiniLM-L6-v2`)
- Dimensionality reduction: UMAP
- Clustering: HDBSCAN
- Tokenization: `sklearn`’s `CountVectorizer`
- Weighing: `c-TF-IDF`
- Representation: `gpt-4.1-nano-2025-04-14` (instead of KeyBERT)

OpenAI embeddings and representation were chosen for their performance in comparison with the free alternatives; the large embeddings model, in particular, was developed with dimensionality reduction in mind.¹⁸ For this dataset size (around 50,000 tokens), the whole analysis cost around a cent.

4.2.1 Interpretation approach

Using the default HDBSCAN `min_cluster_size` of 150 yielded approximately 80 topics, making interpretation challenging (see full dendrogram in Appendix C). After filtering irrelevant topics (e.g., dates, author names) and qualitative classification, four main topics were constructed:

- **AI-generated assets:** content created for media (games, films ads) using AI.
- **Fake or low-quality events:** problematic public events advertised with AI-generated content.
- **Low-quality digital content:** spam, surreal slop, SEO-focused content, ad monetization, and moderation issues.
- **Information and politics:** misinformation, political narratives, media influence, public belief, and the epistemic crisis.

This approach has a limitation: my previous knowledge of the thematic analysis might interfere in the creation of topics. With this in mind, I also create an implementation that is less dependent on manual interpretation.

4.2.2 Filtering approach

Setting `min_cluster_size` to 18 resulted in about 10 topics. Only one was directly related to AI slop, while others included publication dates and named entities (see Figure 6). This step enables us to filter out the irrelevant categories and re-run the model again, now only on the relevant subset (“AI-Generated content and Slop”). Results can be seen in Figure 7.

The resulting topics are very similar to the themes constructed in the thematic analysis. These include ideas of monetization, surreality, spam, advertisement, libraries, SEO, misinformation and Trump, though it has one odd category (“Milton Foundation’s Political Advocacy”).

¹⁸<https://openai.com/index/new-embedding-models-and-api-updates/>




Figure 6: Visualization of topic modelling results before filtering. The plot shows a large cluster of "AI-generated content and slop" items, alongside smaller clusters of less relevant categories.




Figure 7: Results of topic modelling after filtering.

4.3 LLM-based thematic modelling

Recent literature suggests that using off-the-shelf LLMs for classification can be highly effective, even outperforming traditional methods (Chen & Shu, 2024). We test this approach here.

The LLM used was **Gemini 2.5-Pro** through the Google AI Studio interface. At the time of writing, **Gemini-2.5-Pro-Preview-06-05** is the highest placed model on the HuggingFace Chatbot Arena Leaderboard.¹⁹ Standard configurations were used: temperature set to 1, no system instructions and a simple prompt. A second prompt was used to summarize the previous response. This was done only once—no other attempts were made. Both prompts and responses are reproduced in Appendix D. The resulting themes were:

- **Defining Traits - Low Quality and Uncanny:** Slop is defined by its poor quality, visual flaws, nonsensical text, and a “soulless” or cheap feel that betrays its non-human origin.
- **Motivations - Profit, Speed, and Scale:** The creation of slop is driven by economic incentives such as engagement farming, corporate cost-cutting.
- **Impact and Consequences - Deception and Real-World Harm:** Slop has tangible negative effects, such as fake events, misinformation, and political propaganda.
- **The ‘Slop’ Framework - A Polluted Information Ecosystem:** Slop is a symptom of a “post-truth” sociotechnical environment where digital pollution erodes public trust and blurs the lines between truth and fiction.

¹⁹<https://huggingface.co/spaces/lmarena-ai/chatbot-arena-leaderboard>

4.4 Discussion

We compare four approaches: manual thematic analysis, traditional topic modelling using TF-IDF, contextual embedding-based topic modelling, and an LLM-based approach which we refer to as “thematic modelling”.

In summary, the themes identified by the LLM-based method were: traits of low quality and uncanniness; motivations rooted in profit, speed, and scale; consequences such as misinformation and propaganda; and the cultural context of pollution in a post-truth information ecosystem. The topic modelling approach raises similar themes, as well as touching on more concrete and specific issues, such as AI-generated books, Italian Brainrot and AI-driven search results, as well as completely different topics such as ‘Food culture’ and ‘Milton Foundation’.²⁰

The LLM-based approach outperformed the topic modelling methods. It provided nuanced, interpretable outcomes with explanations, avoiding the pitfalls of simpler models that focus on named entities or tangential events. It produced more relevant, abstracted topics (e.g., ‘Defining Traits - Low Quality and Uncanny’) while the topic modelling produced more concrete topics (e.g., ‘Viral Italian Brainrot and Brands’). While the labels in the topic modelling were also LLM-generated, there is considerable benefit in using LLMs directly for this task. Moreover, it was free, fast, and easy to implement. While BERTopic produced strong results, its usage requires specialized knowledge and setting up its pipeline without the dedicated library would be complex, time-consuming, and error-prone.

One benefit of topic modelling, however, is experimentation. While not implemented here, BERTopic supports extensive modification of processing steps, which can significantly alter the quality of the results. In comparison, experimentation with the LLM-based approach is less straightforward, as prompt engineering is more of an art than a science. It relies on models which might change behaviour unexpectedly and requires empirical validation which can become cumbersome.

The LLM-generated themes are remarkably similar to my detailed and painstakingly coded thematic analysis. However, while the generated themes are accurate, they lack the richness of the thematic analysis. For instance, it emphasized textual gibberish in AI-generated images—a relevant feature of earlier datasets, but now technologically outdated. The model’s output remained grounded in descriptive surface patterns, whereas my manual analysis was more interpretive and, arguably, more meaningful. However, the convergence of the results between all methods lends mutual validation and strengthens confidence in the findings.

²⁰While the latter two might speak to some hidden insight, I struggle to interpret them as anything but noise.

4.4.1 Note on LLM-based methods

For exploratory tasks like this, LLMs offer an efficient alternative to traditional methods. In a single five-second query I produced results comparable to a five-hour quantitative analysis and a five-day qualitative analysis. For this small dataset, the generated topics were cheaper, easier, more interpretable, and more nuanced than traditional topic modelling—though this balance might shift with dataset size.

Although efficient, LLM-based methods have limitations. While there are short-term benefits for the analysis at hand, it might result in long-term institutional harms through reducing a researcher’s engagement with their subject matter. This was something I experienced during this work.

While the LLM’s strict adherence to the provided instruction can be a desirable trait, it lacks the contextual understanding of a human researcher. A human might question a flawed premise, whereas the LLM will always oblige. Its output should be treated as a classification that requires human interpretation (much like traditional topic modelling). Care should also be taken with the terminology; these are not the “LLM’s interpretations”, but the results of a “LLM-based method”.

While LLMs cannot replace the subtlety of a thorough thematic analysis, they can be strategically integrated into the research process. Using them as a preliminary step is efficient but risks creating an anchoring effect that biases subsequent analysis. Alternatively, as in this study, they can be used *a posteriori* to validate findings. Unless there is a particular reason to avoid it, incorporating this LLM validation step after a qualitative analysis might be beneficial in most cases.

5 Linguistic and statistical features

This section investigates linguistic and statistical features of AI slop. The analysis is limited to text itself, excluding metadata, fact-checking, and “white box” access to models (Bondielli & Marcelloni, 2019; Wei et al., 2024; Zubiaga et al., 2019).

5.1 SlopNews: a dataset of AI slop

I introduce *SlopNews*, an English-language dataset containing 20,000 slop news articles, 20,000 human-written news articles, and 20,000 machine-generated news articles. The data are not shared directly due to copyright but is easily reproducible.

This dataset focuses on “slop news” over other text-based slop (e.g., submission slop from e-books or academic papers) due to data availability. Ideally, the dataset would differentiate a) slop news from fake news, and b) high- from low-quality AI-generated content. However, these categories were not included in the dataset due to practical limitations.

Articles were labelled as “slop” based on their source website, a methodology also used in tackling fake news (Cameron, 2016). This work intends to enable future classification without relying on such metadata. By using content already identified as slop by others, this research avoids the ethical concerns of applying this label ourselves, such as the potential for reputational harm (e.g., Knibbs, 2024b).

Sources were chosen based on reports from NewsGuard (Sadeghi & Arvanitis, 2023) and DoubleVerify (Luu, 2025),²¹ who identified websites containing generic, AI-generated, or fraudulent content. Four of their examples were selected and scraped: *famadillo*, *filthylucre*, *nbc sport*, and *cbsnews2*. These were selected based on their availability and because they were easier to scrape, creating a convenience sample. As a comparison baseline, we used human-written and machine-generated articles from three datasets:

- NeuralNews : human New York Times articles and AI-generated (GROVER) text (Tan et al., 2020);
- GlobalNews: Kaggle dataset obtained from NewsAPI (Saksham, 2023);²²
- NewsData.io: news data aggregation company which offers free small datasets;²³

All data were pre-processed to remove empty entries, duplicates, and non-English articles (using `langid`), and the non-slop categories were downsampled to create three 20,000-entry samples (slop, human, and AI-generated), for a final total of 60,000 articles. The final composition of the dataset is shown in Table 6.

²¹Others can also be found on Knibbs (2024a). See Appendix E for all websites.

²²<https://newsapi.org/>

²³<https://newsdata.io/datasets>

Table 6: SlopNews data sources.

Source	Amount	Category
NeuralNews	20,000	AI
NeuralNews	8,251	Human
GlobalNews	8,251	Human
Newsdata.io	3,498	Human
nbc sport.com	15,948	Slop
filthy lucre.om	2,608	Slop
famadillo.com	1,136	Slop
cbsnews2.com	715	Slop

5.2 Analysis

The following metrics were computed for each article:

- **Sentence length:** mean, median, and standard deviation of sentence length per document.
- **Token length:** mean, median, and standard deviation of tokens (words/subwords) per document.
- **Vocabulary size:** proportion of unique tokens normalized by document size.
- **Stop word usage:** proportion of stop words normalized by document size.
- **Readability:** Flesch Reading Ease score, which based on sentence length and syllables per word.
- **Complexity:** mean and standard deviation of dependency distance, which is a measure of the distance between related words.
- **Coherence:** cosine similarity between adjacent (first- and second-order) sentence embeddings.
- **Polarity:** sentiment measurer ranging from -1 (negative) to 1 (positive). Calculated using a Naïve Bayes classifier from the `textblob` library.
- **Subjectivity:** ranges from 0 (objective) to 1 (subjective). Calculated using the same classifier as polarity.
- **Sentiment (neural):** computed using outputs from a more sophisticated sentiment analysis model (a deep learning classifier from the `flair` library). Returns a sentence-level binary classification (positive/negative)

and a confidence value (0–1); these were transformed into a document-level score by averaging the sentence values weighed by their confidence.

All metrics were calculated using the `textdescriptives` library unless stated otherwise. Note that several metrics (sentence length, readability, complexity, polarity) are sentence-level measures—creating a distribution for each document.

5.3 Results

Analysis was performed on a sample of the dataset comprised of 19,912 slop and 11,874 non-slop articles (7,975 globalnews, 3,386 newsdataio, 513 neural-news), after excluding all AI-generated articles and all documents with less than 100 characters (<1% of remaining documents) due to parsing errors in some libraries. This also prevented the computation of perplexity measures and introduced some unexplainable selection bias.

Figure 8 shows broadly similar distributions between slop and non-slop, with a few exceptions: slop shows a slightly lower mean in standard deviation of sentence length and a higher concentration of low unique token proportions. Figure 9 shows more expressive differences: slop exhibits a more uniform distribution than non-slop in both sentiment polarity and subjectivity, and a distinctly different distribution in neural sentiment. Neural sentiment for slop is concentrated around positive sentiment, whereas the non-slop distribution has three peaks around negative, neutral, and positive sentiment.²⁴

Statistical comparisons were conducted using two-sided Mann-Whitney U tests, as all distributions are non-normal.²⁵ Since we expect results to be statistically discernible due to the large sample size, we assess its practical implications by computing the Common Language Effect Size (CLES), a measure developed to be intuitive to non-statisticians. CLES indicates the probability that a randomly chosen item from one group scores higher than a randomly chosen item from another. It ranges from zero to one, where 0.5 indicates there is no effect. All measures were computed using the `pingouin` library.

As indicated in Table 7, all metrics show statistically discernible differences ($p < 0.05$), but most effect sizes were negligible. Metrics such as mean sentence length, token length, readability, coherence, and stop word proportion show minimal effect sizes (CLES 0.49–0.51). More substantial (although still modest) differences appear in text complexity and variation: sentence length std. dev. (CLES 0.59), dependency distance (CLES 0.45), dependency distance std. dev. (CLES 0.42), and unique token proportion (CLES 0.40). These may be acting as proxies for perplexity and burstiness, which are standard entropy measures used in AI-generated text detection (Wu et al., 2025). Our findings align with

²⁴Neutral sentiment, in this case, represents low confidence in the classification.

²⁵It has since come to my attention that the test I used to check normality (`pg.normality`) is also sensitive to large samples, which resulted in the diagnosis that no distributions were normal—which is visibly false on the plots. The use of the Mann-Whitney U test shouldn’t be a problem, however, since it can also be used for the comparison of normal distributions.

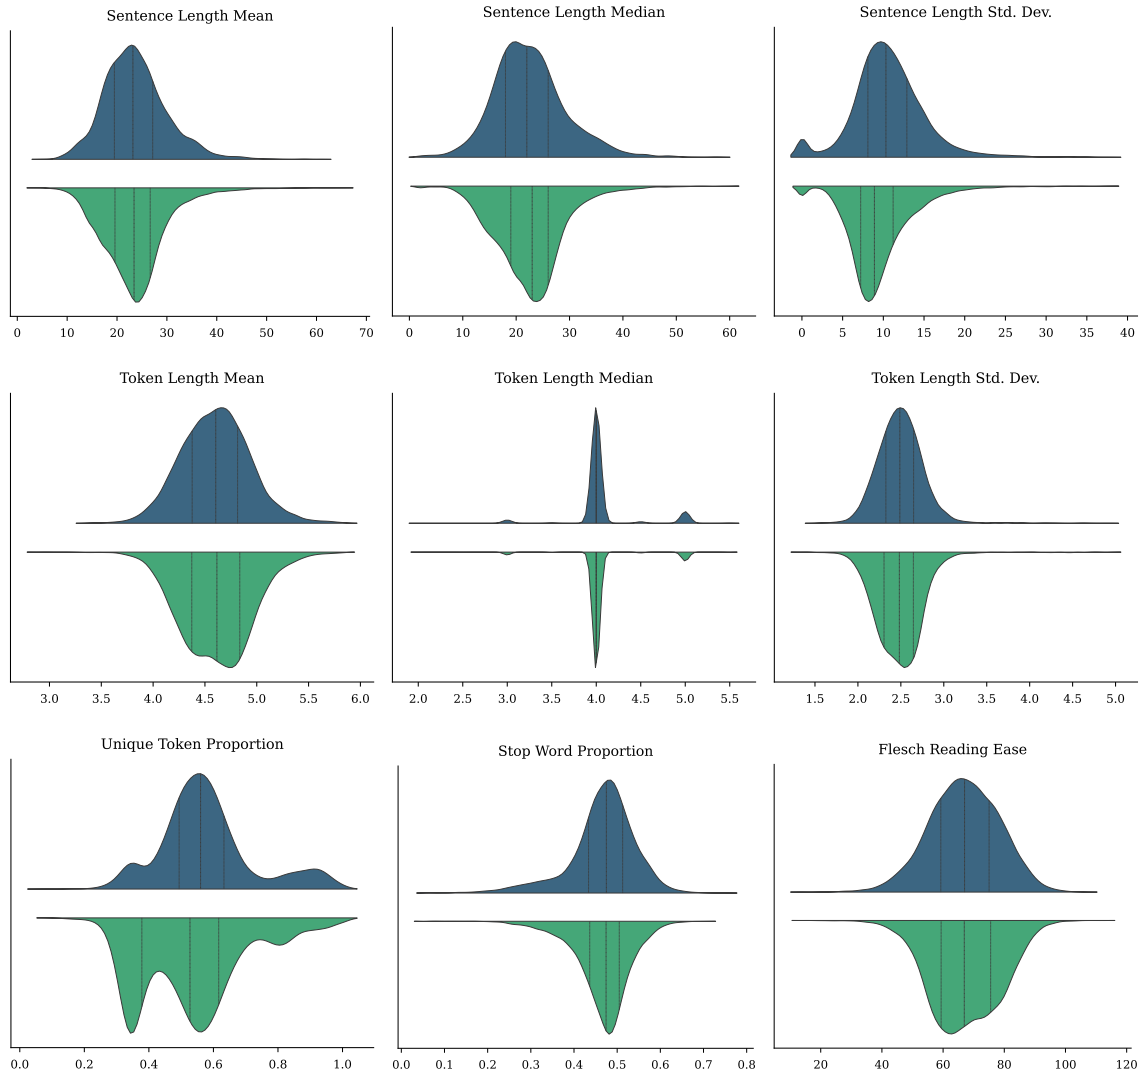


Figure 8: Split violin plots comparing non-slop (top, blue) and slop (bottom, green) conditions across sentence, token, and readability metrics. While most distributions look similar, slop shows a lowered mean in sentence length standard deviation, and a considerable distribution difference in unique token proportions. Outliers (2.3% of the dataset) were discarded for better visualization.

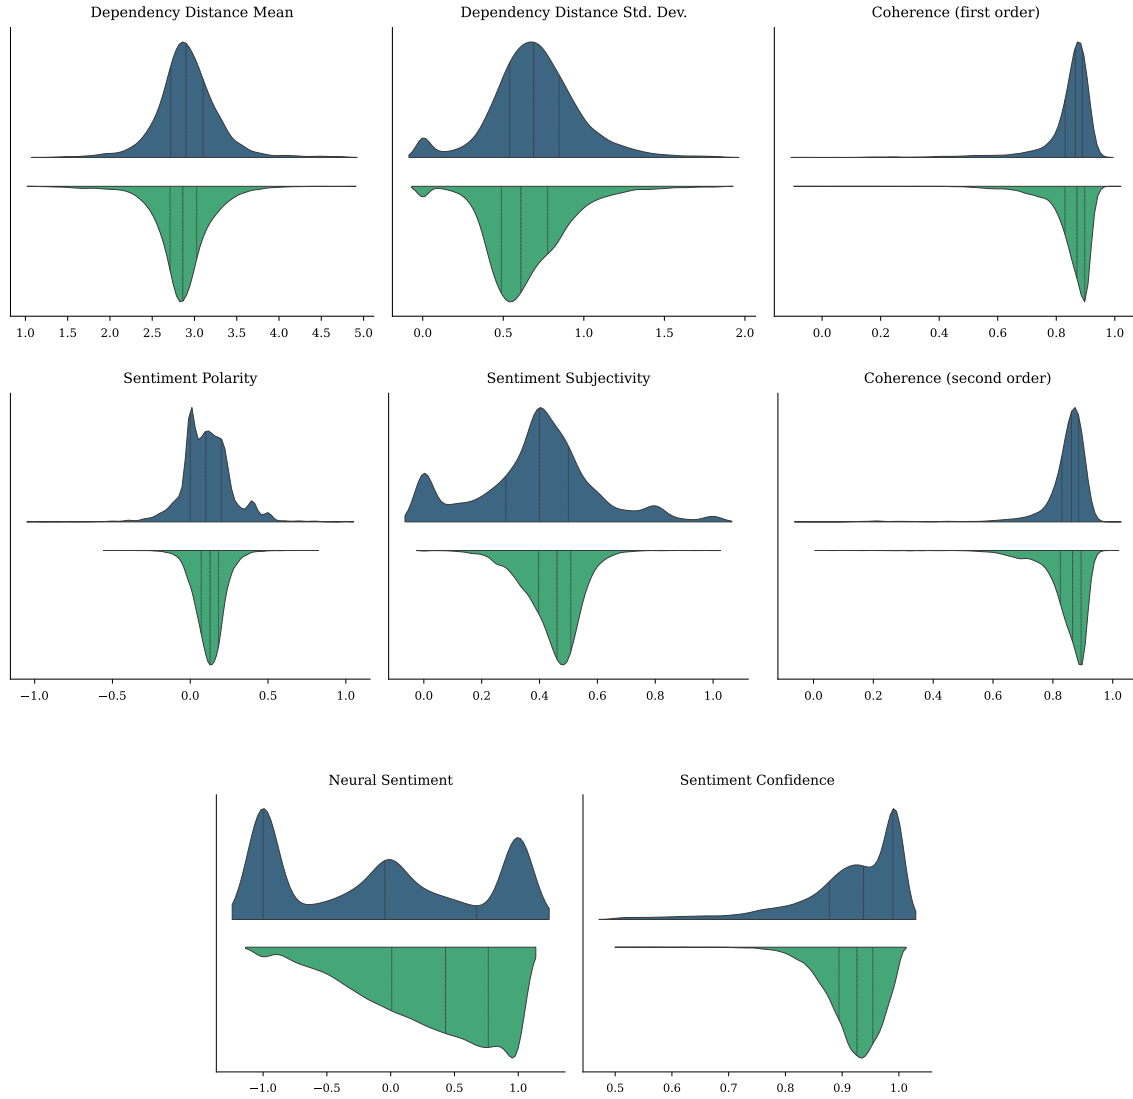


Figure 9: Split violin plots comparing non-slop (top, blue) and slop (bottom, green) conditions across complexity, coherence and sentiment measures. Slop displays less variation in all sentiment measures. Outliers (2.3% of the dataset) were discarded for better visualization.

Table 7: Results of Mann–Whitney U-tests comparing non-slop and slop samples. CLES = Common Language Effect Size. Values truncated for brevity. Larger effect sizes are bolded.

Metric	U	p	Non-slop Mean	Slop Mean	CLES
Sentence Length Mean	118668692.0	0.001	24.380	24.796	0.51
Sentence Length Median	114212087.5	0.008	23.135	23.802	0.49
Sentence Length Std. Dev.	139240493.5	3e-190	10.894	9.989	0.59
Token Length Mean	121076067.0	3-e4	4.616	4.633	0.51
Token Length Median	116173757.0	4-e6	4.080	4.064	0.49
Token Length Std. Dev.	116222253.0	0.011	2.544	2.544	0.49
Flesch Reading Ease	120226097.5	0.011	65.987	65.774	0.50
Dependency Distance Mean	108375684.0	1e-35	2.948	2.877	0.45
Dependency Distance Std. Dev.	100868767.5	1e-106	0.709	0.654	0.42
Coherence (first order)	119395882.5	5e-24	0.837	0.847	0.53
Coherence (second order)	106350805.5	6e-11	0.841	0.844	0.52
Unique Token Proportion	95902758.0	6e-175	0.581	0.529	0.40
Stop Word Proportion	54229874.0	0.018	0.463	0.464	0.49
Sentiment Polarity	130650555.0	1e-55	0.111	0.125	0.55
Sentiment Subjectivity	146451786.5	8e-279	0.386	0.448	0.61
Neural Sentiment	144980342.5	0.0	-0.089	0.351	0.65
Sentiment Confidence	97631371.5	3e-70	0.913	0.920	0.43

the literature which shows that AI text tends to be more uniform than human-written text (Chen & Shu, 2024). We thus conclude that **slop text has lower variation and complexity than non-slop text**.

Slop articles also show higher sentiment scores in both simple (CLES 0.55) and neural (CLES 0.65) models, as well as greater subjectivity (CLES 0.61). These findings suggest that **slop tends to exhibit more positive sentiment and subjectivity than non-slop**. One possible explanation is that slop news aims to be unobtrusive, adopting a generic positivity. This contrasts with other kinds of slop, such as social media slop and disinformation, where the purpose is being attention-grabbing.

In summary, slop text tends to be slightly less varied and less complex, while also being more positive and subjective. These differences are measurable but modest. Of the features investigated, lexical diversity, dependency structure, and sentiment offer the most promise for slop classification.

5.4 Limitations and future work

Similar to misinformation (Dufour et al., 2024), AI slop is primarily image-based. This study deliberately focuses on text. Even so, the dataset used here is quite limited, as reflected in the selection bias incurred from successive convenience sampling. Future research can be considerably improved by incorporating more diverse data sources, other types of slop, and languages other than English. Furthermore, a wider range of linguistic metrics can be applied, such as:

- Known slop fingerprints, such as accidental prompt disclosure, model refusals, vocabulary patterns, emoji usage and markup features. These can be sourced from efforts such as the Wikipedia Project AI Cleanup²⁶, the Problematic Paper Screener²⁷, originality.ai blog (Guillham, 2024), and Pangram Resources for Educators²⁸.
- Valence, arousal and dominance (VAD) metrics (Mohammad, 2025; Vishnubhotla & Mohammad, 2022)
- Linguistic Inquiry and Word Count (LIWC) features²⁹
- Givenness as a measure of complexity³⁰
- Type Token Ratios (TTR) as a measure of lexical diversity³¹

It is important to emphasize that the statistical findings presented here are unlikely to be generalizable. The observed differences are relatively weak and likely reflect biases (e.g. stylistic traits) in the data source selection rather than differences intrinsic to the slop/non-slop distinction. While these metrics reveal

²⁶https://en.wikipedia.org/wiki/Wikipedia:WikiProject_AI_Cleanup/AI_catchphrases

²⁷<https://www.irit.fr/~Guillaume.Cabanac/problematic-paper-screener>

²⁸<https://www.pangram.com/resources/most-common-ai-phrases>

²⁹https://lit.eecs.umich.edu/geoliwc/liwc_dictionary.html

³⁰https://trunajod20.readthedocs.io/en/latest/api_reference/givenness.html

³¹https://trunajod20.readthedocs.io/en/latest/api_reference/ttr.html

some distributional differences, they are not robust enough to support reliable classification or decision-making on their own.

A natural follow-up for this work would be the development of a machine-learning classification model incorporating the features identified here. If its performance is satisfactory, it can be published as an open-source library or API to be used in news aggregation, moderation, fact-checking, etc.

Finally, future work could aim to operationalize the characteristics investigated in the first part of this study (e.g., writing quality; Chakrabarty et al., 2025). However, this might prove difficult, as many are highly subjective and rely as much on the reader’s interpretation as on the content itself.

6 Conclusion

This work investigated AI slop—its qualities, metaphors, themes, types, and statistical characteristics. Slop is described as excessive, inaccurate, inauthentic, mediocre, surreal, distorted, generic, pointless, and/or kitschy. It is enabled by cheap, scalable automation which is used to respond to the incentives of the engagement economy. AI slop floods digital spaces with low-effort, contextless, generic and/or surreal material. It is created for productivity, profit, or manipulation rather than meaningful communication. It is created by corporations optimizing for scale and individuals chasing side income. While present in text, images, audio and video, slop images are its most iconic form, as illustrated by Shrimp Jesus and other “weird AI crap” (Read, 2024).

AI slop operates at the limits of moderation, as it is not necessarily factually incorrect (like fake news) or offensive (like hate speech). Yet, its informational effects can be understood in the context of “pollution of information ecosystem and loss of consensus reality” (Slattery et al., 2025). By its sheer scale, it reduces the signal-to-noise ratio of the information environment and demands additional effort in order to filter it.

Discourse about slop can be categorized into four main themes: efficiency and productivity (slop is cheap, easy and scalable), danger and discernment (slop is threatening and requires filtering), epistemic crisis (slop is related to post-truth) and engagement economy (slop is created because of platform incentives). It is often described metaphorically in terms of flooding, contamination, infestation and rubbish.

The four main types of AI slop are: social media slop (low-effort, outrageous or manipulative engagement bait created by “sloppers” for the purpose of profiting off social media algorithms), corporate slop (low-quality AI-generated media used by corporations for cost-cutting purposes, typically in advertising), submission slop (papers, books, job applications, and other media which are mass-submitted and overwhelm selection systems) and slop news (low-quality news articles created by professional sloppers for SEO manipulation).

By analysing its linguistic and statistical characteristics, we find that slop news is less complex, less varied and more positive in comparison with non-slop news. The most noticeable difference occurs in the distribution of sentiment across slop and non-slop, as there are almost no slop news articles with negative sentiment.

This study is, to the best of my knowledge, the first systematic analysis of AI slop. Although limited, it outlines potential directions for further research and lays the groundwork for developing interpretable systems for slop detection. By advancing our understanding of AI slop, this work aims to support efforts to safeguard the integrity of the information environment.

Word count

This work contains around 10,000 words (excluding supplementary materials).

AI Use Declaration

I have used the following AI tools for the following reasons:

- **ChatGPT and Gemini:** text editing and typesetting with the purpose of improving delivery, clarity and information density.
- **Cursor:** scraping and data visualization boilerplate

I declare that this submission is my own original work and that any AI-generated text or images have been fully cited as such.

This document was typeset using L^AT_EX. The cover art is sourced from https://commons.wikimedia.org/wiki/File:Shrimp_Jesus_example.jpg.

References

- Aich, A., Bhattacharya, S., & Parde, N. (2022). Demystifying Neural Fake News via Linguistic Feature-Based Interpretation. *Proceedings of the 29th International Conference on Computational Linguistics*, 6586–6599.
- Aleksic, A. (2024, November 12). *The linguistic observer effect*. The Etymology Nerd. Retrieved May 13, 2025, from <https://etymology.substack.com/p/the-linguistic-observer-effect>
- Allen, J. D., Howland, B., Mobius, M., Rothschild, D., & Watts, D. J. (2020). Evaluating the fake news problem at the scale of the information ecosystem. *Science Advances*, 6(14). <https://doi.org/10.1126/sciadv.aay3539>
- Arguedas, A. R., Robertson, C. T., Fletcher, R., & Nielsen, R. K. (2022). *Echo chambers, filter bubbles, and polarisation: A literature review*. Reuters Institute for the Study of Journalism. <https://doi.org/10.60625/RISJ-ETXJ-7K60>
- Bailey, R. (2024, May 4). *The case of the AI-generated giant rat penis*. Reason.com. Retrieved July 4, 2025, from <https://reason.com/2024/05/04/the-case-of-the-ai-generated-giant-rat-penis/>
- Barker, E. (2011). *Stepping out of the Ivory Tower: A Sociological Engagement in 'The Cult Wars'*. Retrieved May 30, 2025, from <https://journals.sagepub.com/doi/abs/10.4256/mio.2010.0026>
- Barthel, M., Mitchell, A., & Holcomb, J. (2016, December 15). *Many Americans Believe Fake News Is Sowing Confusion*. Pew Research Center. Retrieved May 7, 2025, from <https://www.pewresearch.org/journalism/2016/12/15/many-americans-believe-fake-news-is-sowing-confusion/>
- Bik, E. (2024, February 15). *The rat with the big balls and the enormous penis – how Frontiers published a paper with botched AI-generated images*. Science Integrity Digest. Retrieved July 8, 2025, from <https://scienceintegritydigest.com/2024/02/15/the-rat-with-the-big-balls-and-enormous-penis-how-frontiers-published-a-paper-with-botched-ai-generated-images/>
- Bondielli, A., & Marcelloni, F. (2019). A survey on fake news and rumour detection techniques. *Information Sciences*, 497, 38–55. <https://doi.org/10.1016/j.ins.2019.05.035>
- Braun, V., & Clarke, V. (2006). Using thematic analysis in psychology. *Qualitative Research in Psychology*, 3(2), 77–101. <https://doi.org/10.1191/1478088706qp063oa>
- Cameron, D. (2016, November 18). *Viral 'fake news' list deleted after author and students get harassed, doxed*. The Daily Dot. Retrieved May 26, 2025, from <https://www.dailydot.com/debug/fake-news-facebook-deleted-alt-right/>
- Carpenter, J. (2025, March 4). *If not friend, why friend-shaped?* Retrieved May 13, 2025, from <https://jgcarpenter.com/blog.html?blogPost=if-not-friend-why-friend-shaped-julie-carpenter>

- Cellan-Jones, R. (2017). Fake news worries 'are growing' suggests BBC poll [newspaper]. *BBC News: Technology*. Retrieved May 29, 2025, from <https://www.bbc.com/news/technology-41319683>
- Chakrabarty, T., Laban, P., & Wu, C.-S. (2025). *AI-Slop to AI-Polish? Aligning Language Models through Edit-Based Writing Rewards and Test-time Computation*. arXiv: 2504.07532 [cs]. <https://doi.org/10.48550/arXiv.2504.07532>
- Chen, C., & Shu, K. (2024). Combating misinformation in the age of LLMs: Opportunities and challenges. *AI Magazine*, 45(3), 354–368. <https://doi.org/10.1002/aaai.12188>
- Chkhaidze, A., Buyruk, P., & Boroditsky, L. (2021). Linguistic Metaphors Shape Attitudes towards Immigration. *Proceedings of the Annual Meeting of the Cognitive Science Society*, 43(43). Retrieved June 20, 2025, from <https://escholarship.org/uc/item/1xf1t3vv>
- Clarke, N. (2023, December 2). *A Concerning Trend*. Retrieved May 28, 2025, from <https://neil-clarke.com/a-concerning-trend/>
- Cohen, G. A. (2002, June 6). Deeper into Bullshit. In S. Buss & L. Overton (Eds.), *The Contours of Agency* (pp. 321–339). The MIT Press. <https://doi.org/10.7551/mitpress/2143.003.0015>
- Costello, T., Pennycook, G., & Rand, D. (2025). *Just the facts: How dialogues with AI reduce conspiracy beliefs*. <https://doi.org/10.31234/osf.io/h7n8u.v1>
- Dalton, C. (2016). Bullshit for you; transcendence for me. A commentary on “On the reception and detection of pseudo-profound bullshit”. *Judgment and Decision Making*, 11(1), 121–122. <https://doi.org/10.1017/S1930297500007646>
- Derrida, J. (2020). *A Farmácia de Platão*. Iluminuras.
- Doctorow, C. (2023). The ‘Enshittification’ of TikTok [magazine]. *Wired*. Retrieved May 27, 2025, from <https://www.wired.com/story/tiktok-platforms-cory-doctorow/>
- Doughman, J., Afzal, O. M., Toyin, H. O., Shehata, S., Nakov, P., & Talat, Z. (2024). *Exploring the Limitations of Detecting Machine-Generated Text*. arXiv: 2406.11073 [cs]. <https://doi.org/10.48550/arXiv.2406.11073>
- Dufour, N., Pathak, A., Samangouei, P., Hariri, N., Deshetti, S., Dudfield, A., Guess, C., Escayola, P. H., Tran, B., Babakar, M., & Bregler, C. (2024, May 21). *AMMeBa: A Large-Scale Survey and Dataset of Media-Based Misinformation In-The-Wild* (2). arXiv: 2405.11697 [cs]. <https://doi.org/10.48550/arXiv.2405.11697>
- Easwaran, K. (2023). Bullshit activities. *Analytic Philosophy*, 1–23. <https://doi.org/10.1111/phib.12328>
- Emi, B., & Spero, M. (2024). *Technical Report on the Pangram AI-Generated Text Classifier*. arXiv: 2402.14873 [cs]. <https://doi.org/10.48550/arXiv.2402.14873>
- Flanagan, D., Mavergames, C., O’Brien, A., Parker, S., Razis, G., & Zhou, H. (2024). *Upholding Research Integrity in the Age of AI*. Wiley. Retrieved

- June 30, 2025, from https://res6.info.wiley.com/res/tracking/Wiley_State_of_the_Future_2024-Digital_Transformation.pdf
- Forlini, E. (2023, July 25). *OpenAI Quietly Shuts Down AI Text-Detection Tool Over Inaccuracies*. PCMAG. Retrieved July 4, 2025, from <https://www.pcmag.com/news/openai-quietly-shuts-down-ai-text-detection-tool-over-inaccuracies>
- Frankfurt, H. G. (2009). *On Bullshit*. Princeton University Press. <https://doi.org/10.1515/9781400826537>
- Gonen, H., & Goldberg, Y. (2019). *Lipstick on a Pig: Debiasing Methods Cover up Systematic Gender Biases in Word Embeddings But do not Remove Them*. arXiv.org. Retrieved March 11, 2025, from <https://arxiv.org/abs/1903.03862v2>
- Gorrieri, L. (2024). Is ChatGPT Full of Bullshit? *Journal of Ethics and Emerging Technologies*, 34(1), 1–16. <https://doi.org/10.55613/jeet.v34i1.149>
- Graña, R. B. (2016, May 9). *Heidegger ou as Vicissitudes da Destruição*. Age.
- Guillham, J. (2024, August 8). *Most Commonly Used ChatGPT Words & Phrases*. Originality.AI. Retrieved June 30, 2025, from <https://originality.ai/blog/can-humans-detect-chatgpt>
- Han, B.-C. (2015). *The burnout society*. Stanford University Press.
- Haraway, D. (2015). *Simians, cyborgs, and women: The reinvention of nature*. Routledge.
- Heap, T., Lawson, T., Farnik, L., & Aitchison, L. (2025). *Sparse Autoencoders Can Interpret Randomly Initialized Transformers*. arXiv: 2501.17727 [cs]. <https://doi.org/10.48550/arXiv.2501.17727>
- Hicks, M. T., Humphries, J., & Slater, J. (2024). ChatGPT is bullshit. *Ethics and Information Technology*, 26(2), 38. <https://doi.org/10.1007/s10676-024-09775-5>
- Hoffman, B. (2024). Is Slop A.I.’s Answer to Spam? A Phrase Emerges for Bad Search [newspaper]. *The New York Times*. Retrieved May 13, 2025, from <https://www.nytimes.com/2024/06/11/style/ai-search-slop.html>
- Hoffmann, M., & Frase, H. (2023). *Adding Structure to AI Harm: An Introduction to CSET’s AI Harm Framework*. Center for Security and Emerging Technology. Retrieved July 7, 2025, from <https://cset.georgetown.edu/publication/adding-structure-to-ai-harm/>
- Hoover, A. (2025, April 2). *Cheating on tech interviews is soaring. Managers don’t know what to do*. Business Insider. Retrieved May 29, 2025, from <https://www.businessinsider.com/cheating-tech-interviews-soaring-managers-lost-gen-ai-chatgpt-coding-2025-4>
- Huang, L., Yu, W., Ma, W., Zhong, W., Feng, Z., Wang, H., Chen, Q., Peng, W., Feng, X., Qin, B., & Liu, T. (2025). A Survey on Hallucination in Large Language Models: Principles, Taxonomy, Challenges, and Open Questions. *ACM Trans. Inf. Syst.*, 43(2), 42:1–42:55. <https://doi.org/10.1145/3703155>
- Illing, S. (2020, January 16). *“Flood the zone with shit”: How misinformation overwhelmed our democracy*. Vox. Retrieved May 30, 2025, from

- <https://www.vox.com/policy-and-politics/2020/1/16/20991816/impeachment-trump-bannon-misinformation>
- Infante, G., Lecouteux, G., & Sugden, R. (2016). Preference Purification and the Inner Rational Agent: A Critique of the Conventional Wisdom of Behavioural Welfare Economics. *Journal of Economic Methodology*, 23(1), 1–25. <https://doi.org/10.1080/1350178x.2015.1070527>
- Joelving, F. (2025, March 4). *Was nonsense ‘vegetative electron microscopy’ phrase a Farsi typo?* Retraction Watch. Retrieved May 23, 2025, from <https://retractionwatch.com/2025/03/04/vegetative-electron-microscopy-phrase-farsi-typo/>
- Kapoor, S., Henderson, P., & Narayanan, A. (2024). *Promises and pitfalls of artificial intelligence for legal applications*. arXiv: 2402.01656 [cs]. <https://doi.org/10.48550/arXiv.2402.01656>
- King, T., Wu, Z., Koshiyama, A., Kazim, E., & Treleaven, P. (2024). *HEARTS: A Holistic Framework for Explainable, Sustainable and Robust Text Stereotype Detection*. arXiv: 2409.11579. <https://doi.org/10.48550/arXiv.2409.11579>
- Klincewicz, M., Alfano, M., & Fard, A. E. (2025). *Slopaganda: The interaction between propaganda and generative AI*. arXiv: 2503.01560 [cs]. <https://doi.org/10.48550/arXiv.2503.01560>
- Knibbs, K. (2024a). Confessions of an AI Clickbait Kingpin [magazine]. *Wired*. Retrieved May 13, 2025, from <https://www.wired.com/story/confessions-of-an-ai-clickbait-kingpin/>
- Knibbs, K. (2024b). The Guy Behind the Fake AI Halloween Parade Listing Says You’ve Got It All Wrong [magazine]. *Wired*. Retrieved June 29, 2025, from <https://www.wired.com/story/ai-halloween-parade-listing-dublin-interview/>
- Koebler, J. (2025, April 28). *Researchers Secretly Ran a Massive, Unauthorized AI Persuasion Experiment on Reddit Users*. Retrieved May 7, 2025, from <https://www.404media.co/researchers-secretly-ran-a-massive-unauthorized-ai-persuasion-experiment-on-reddit-users/>
- Krenak, A. (2020). *Ideas to postpone the end of the world* (A. Doyle, Trans.). Anansi International.
- Labbé, C., Joelving, F., & Cabanac, G. (2025a, January 29). *Fake papers are contaminating the world’s scientific literature, fueling a corrupt industry and slowing legitimate lifesaving medical research*. The Conversation. Retrieved May 26, 2025, from <http://theconversation.com/fake-papers-are-contaminating-the-worlds-scientific-literature-fueling-a-corrupt-industry-and-slowing-legitimate-lifesaving-medical-research-246224>
- Labbé, C., Joelving, F., & Cabanac, G. (2025b, January 29). *Problematic Paper Screener: Trawling for fraud in the scientific literature*. The Conversation. Retrieved May 28, 2025, from <http://theconversation.com/problematic-paper-screener-trawling-for-fraud-in-the-scientific-literature-246317>
- Lakoff, G., & Johnson, M. (2008). *Metaphors We Live By*. University of Chicago Press.

- Lazer, D., Baum, M. A., Benkler, Y., Berinsky, A. J., Greenhill, K. M., Menczer, F., Metzger, M. J., Nyhan, B., Rand, D. G., Pennycook, G., Rothschild, D., Schudson, M., Sloman, S. A., Cass R. Sunstein, Sunstein, C. R., Thorson, E. A., Duncan J. Watts, Watts, D. J., & Zittrain, J. L. (2018). The science of fake news. *Science*, 359(6380), 1094–1096. <https://doi.org/10.1126/science.aao2998>
MAG ID: 2791544114.
- LowercaseJai (**typedirector**). (2025, June 3). *An Autopsy of House M.D.* [Video]. Retrieved July 1, 2025, from <https://www.youtube.com/watch?v=K7jUw0GCWck>
- Luu, K. (2025, January 15). *The Rise of AI-Generated Slop Sites: What Advertisers Need to Know*. DoubleVerify. Retrieved June 29, 2025, from <https://doubleverify.com/the-rise-of-ai-generated-slop-sites-what-advertisers-need-to-know/>
- Lyotard, J.-F. (1984). *The Postmodern Condition: A report on knowledge*. Retrieved June 16, 2025, from <https://manchesteruniversitypress.co.uk/9780719014505/>
- Maiberg, E. (2025, February 4). *AI-Generated Slop Is Already In Your Public Library*. 404 Media. Retrieved June 23, 2025, from <https://www.404media.co/email/7a1109dd-e6cd-46ac-be2f-a6d6461fc2ff/>
- McMahon, L., & Kleinman, Z. (2024, May 24). *Google AI search tells users to glue pizza and eat rocks*. BBC News. Retrieved July 4, 2025, from <https://www.bbc.com/news/articles/cd11gzejgz4o>
- Mendelsohn, J., & Budak, C. (2025). *When People are Floods: Analyzing Dehumanizing Metaphors in Immigration Discourse with Large Language Models*. arXiv: 2502.13246 [cs]. <https://doi.org/10.48550/arXiv.2502.13246>
- Mohammad, S. M. (2025). *NRC VAD Lexicon v2: Norms for Valence, Arousal, and Dominance for over 55k English Terms*. arXiv: 2503.23547 [cs]. <https://doi.org/10.48550/arXiv.2503.23547>
- Mollick, E. (2025, January 26). *Setting time on fire and the temptation of The Button*. Retrieved June 30, 2025, from <https://www.oneusefulthing.org/p/setting-time-on-fire-and-the-temptation>
- Muzumdar, P., Cheemalapati, S., RamiReddy, S. R., Singh, K., Kurian, G., & Muley, A. (2025). The Dead Internet Theory: A Survey on Artificial Interactions and the Future of Social Media. *Asian Journal of Research in Computer Science*, 18(1), 67–73. <https://doi.org/10.9734/ajrcos/2025/v18i1549>
- Narayanan, A., & Kapoor, S. (2025). *AI as Normal Technology*. Knight First Amendment Institute. Retrieved June 9, 2025, from <http://knightcolumbia.org/content/ai-as-normal-technology>
- Oxford Word of the Year 2024*. (n.d.). Oxford University Press. Retrieved July 4, 2025, from <https://corp.oup.com/word-of-the-year/>
- Pennycook, G., Cheyne, J. A., Barr, N., Koehler, D. J., & Fugelsang, J. A. (2015). On the reception and detection of pseudo-profound bullshit.

- Judgment and Decision Making*, 10(6), 549–563. <https://doi.org/10.1017/S1930297500006999>
- Pennycook, G., & Rand, D. G. (2019). Lazy, not biased: Susceptibility to partisan fake news is better explained by lack of reasoning than by motivated reasoning. *Cognition*, 188, 39–50. <https://doi.org/10.1016/j.cognition.2018.06.011>
- Peters, D. (2022, August 9). *How NewsGuard monitors 8,000 websites to combat misinformation*. Newswhip. Retrieved May 26, 2025, from <https://www.newswhip.com/2022/08/how-newsguard-monitors-8000-websites-to-combat-misinformation/>
- Placido, D. D. (2024, April 28). *Facebook’s AI-Generated ‘Shrimp Jesus,’ Explained*. Forbes. Retrieved May 26, 2025, from <https://www.forbes.com/sites/danidiplacido/2024/04/28/facebooks-surreal-shrimp-jesus-trend-explained/>
- Purdy, K. (2025, May 7). *Open source project curl is sick of users submitting “AI slop” vulnerabilities*. Ars Technica. Retrieved May 28, 2025, from <https://arstechnica.com/gadgets/2025/05/open-source-project-curl-is-sick-of-users-submitting-ai-slop-vulnerabilities/>
- Read, M. (2024, September 25). *Drowning in Slop*. Intelligencer. Retrieved June 30, 2025, from <https://nymag.com/intelligencer/article/ai-generated-content-internet-online-slop-spam.html>
- Roberts, D. (2019, November 16). *With impeachment, America’s epistemic crisis has arrived*. Vox. Retrieved July 1, 2025, from <https://www.vox.com/policy-and-politics/2019/11/16/20964281/impeachment-hearings-trump-america-epistemic-crisis>
- Sadasivan, V. S., Kumar, A., Balasubramanian, S., Wang, W., & Feizi, S. (2023). *Can AI-Generated Text be Reliably Detected?* arXiv: 2303.11156 [cs]. Retrieved June 9, 2023, from <http://arxiv.org/abs/2303.11156>
- Sadeghi, M., & Arvanitis, L. (2023, May 1). *Rise of the Newsbots: AI-Generated News Websites Proliferating Online*. NewsGuard. Retrieved May 26, 2025, from <https://www.newsguardtech.com/special-reports/newsbots-ai-generated-news-websites-proliferating>
- Saksham, K. (2023). *Global News Dataset*. Kaggle. <https://doi.org/10.34740/KAGGLE/DSV/7105651>
- Saphra, N., & Wiegrefe, S. (2024). *Mechanistic?* arXiv: 2410.09087 [cs]. <https://doi.org/10.48550/arXiv.2410.09087>
- Singh, M. (2024). Don’t Believe What They’re Telling You About Misinformation [magazine]. *The New Yorker*. Retrieved May 29, 2025, from <https://www.newyorker.com/magazine/2024/04/22/dont-believe-what-theyre-telling-you-about-misinformation>
- Slattery, P., Saeri, A. K., Grundy, E. A. C., Graham, J., Noetel, M., Uuk, R., Dao, J., Pour, S., Casper, S., & Thompson, N. (2025). *The AI Risk Repository: A Comprehensive Meta-Review, Database, and Taxonomy of Risks From Artificial Intelligence*. arXiv: 2408.12622 [cs]. <https://doi.org/10.48550/arXiv.2408.12622>

- Sperber, D. (2010). The Guru Effect. *Review of Philosophy and Psychology*, 1(4), 583–592. <https://doi.org/10.1007/s13164-010-0025-0>
- Sperber, D., & Wilson, D. (1987). Précis of Relevance: Communication and Cognition. *Behavioral and Brain Sciences*, 10(4), 697–710. <https://doi.org/10.1017/S0140525X00055345>
- Squires, V. (2023). Thematic Analysis. In J. M. Okoko, S. Tunison, & K. D. Walker (Eds.), *Varieties of Qualitative Research Methods: Selected Contextual Perspectives* (pp. 463–468). Springer International Publishing. https://doi.org/10.1007/978-3-031-04394-9_72
- Stockwell, S. (2024). *AI-Enabled Influence Operations: Threat Analysis of the 2024 UK and European Elections*. Centre for Emerging Technology and Security. Retrieved July 3, 2025, from <https://cetas.turing.ac.uk/publications/ai-enabled-influence-operations-threat-analysis-2024-uk-and-european-elections>
- Sullivan, D., & Nelson, C. (2023, February 8). *Google Search’s guidance about AI-generated content*. Google Search Central Blog. Retrieved June 30, 2025, from <https://developers.google.com/search/blog/2023/02/google-search-and-ai-content>
- Tan, R., Plummer, B. A., & Saenko, K. (2020). Detecting Cross-Modal Inconsistency to Defend Against Neural Fake News. *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2081–2106. <https://doi.org/10.18653/v1/2020.emnlp-main.163>
- Tang, R., Chuang, Y.-N., & Hu, X. (2023). *The Science of Detecting LLM-Generated Texts*. arXiv: 2303.07205 [cs]. <https://doi.org/10.48550/arXiv.2303.07205>
- Tangermann, V. (2024, January 12). *Amazon Is Selling Products With AI-Generated Names Like "I Cannot Fulfill This Request It Goes Against OpenAI Use Policy"*. Futurism. Retrieved May 30, 2025, from <https://futurism.com/amazon-products-ai-generated>
- Templeton, A., Conerly, T., Marcus, J., Lindsey, J., Bricken, T., Chen, B., Adam Pearce, Citro, C., Ameisen, E., Jones, A., Cunningham, H., Nicholas L Turner, McDougall, C., MacDiarmid, M., Tamkin, A., Durmus, E., Hume, T., Mosconi, F., Freeman, C. D., . . . Henighan, T. (2024). *Scaling Monosemanticity: Extracting Interpretable Features from Claude 3 Sonnet*. Anthropic. Retrieved June 30, 2025, from <https://transformer-circuits.pub/2024/scaling-monosemanticity/index.html>
- Tran, T. H. (2023, May 6). *How This Doctor Wrote Dozens of Science Papers With ChatGPT*. The Daily Beast. Retrieved May 28, 2025, from <https://www.thedailybeast.com/how-this-doctor-wrote-dozens-of-science-papers-with-chatgpt/>
- Trevisan, A., Giddens, H., Dillon, S., & Blackwell, A. F. (2024). *Measuring Bullshit in the Language Games played by ChatGPT*. arXiv: 2411.15129 [cs]. <https://doi.org/10.48550/arXiv.2411.15129>

- Tully, S., Longoni, C., & Appel, G. (2025). EXPRESS: Lower Artificial Intelligence Literacy Predicts Greater AI Receptivity. *Journal of Marketing*, 00222429251314491. <https://doi.org/10.1177/00222429251314491>
- Tversky, A., & Kahneman, D. (1974). Judgment under Uncertainty: Heuristics and Biases. *Science*. <https://doi.org/10.1126/science.185.4157.1124>
- Vincent, J. (2023, February 2). *Hustle bros are jumping on the AI bandwagon*. The Verge. Retrieved July 8, 2025, from <https://www.theverge.com/2023/2/2/23582772/chatgpt-ai-get-rich-quick-schemes-hustlers-web>
- Vishnubhotla, K., & Mohammad, S. M. (2022). Tweet Emotion Dynamics: Emotion Word Usage in Tweets from US and Canada. *Proceedings of the Thirteenth Language Resources and Evaluation Conference*, 4162–4176. Retrieved July 2, 2025, from <https://aclanthology.org/2022.lrec-1.442/>
- Wardle, C., & Derakhshan, H. (2017). *Information disorder: Toward an interdisciplinary framework for research and policy making*. Retrieved May 13, 2025, from <https://edoc.coe.int/en/media/7495-information-disorder-toward-an-interdisciplinary-framework-for-research-and-policy-making.html>
- Weber-Wulff, D., Anohina-Naumeca, A., Bjelobaba, S., Foltýnek, T., Guerrero-Dib, J., Popoola, O., Šigut, P., & Waddington, L. (2023). Testing of detection tools for AI-generated text. *International Journal for Educational Integrity*, 19(1), 26. <https://doi.org/10.1007/s40979-023-00146-z>
- Wei, J., Yang, C., Song, X., Lu, Y., Hu, N., Huang, J., Tran, D., Peng, D., Liu, R., Huang, D., Du, C., & Le, Q. V. (2024). *Long-form factuality in large language models*. arXiv: 2403.18802 [cs]. <https://doi.org/10.48550/arXiv.2403.18802>
- Willison, S. (2024, May 8). *Slop is the new name for unwanted AI-generated content*. Simon Willison's Weblog. Retrieved May 13, 2025, from <https://simonwillison.net/2024/May/8/slop/>
- World Economic Forum. (2024). *Global Risks 2024: Disinformation Tops Global Risks 2024 as Environmental Threats Intensify*. World Economic Forum. Retrieved May 30, 2025, from <https://www.weforum.org/press/2024/01/global-risks-report-2024-press-release/>
- Wu, J., Yang, S., Zhan, R., Yuan, Y., Chao, L. S., & Wong, D. F. (2025). A Survey on LLM-Generated Text Detection: Necessity, Methods, and Future Directions. *Computational Linguistics*, 51(1), 275–338. https://doi.org/10.1162/coli.a_00549
- Wynn-Williams, S. (2025). *Careless people: A cautionary tale of power, greed, and lost idealism* (First edition). Flatiron Books.
- Zubiaga, A., Aker, A., Bontcheva, K., Liakata, M., & Procter, R. (2019). Detection and Resolution of Rumours in Social Media: A Survey. *ACM Computing Surveys*, 51(2), 1–36. <https://doi.org/10.1145/3161603>

A Appendix: Articles about slop

#	Name	Source	Topic
1	AI slop	Wikipedia	AI Slop
2	Spam, junk ... slop? The latest wave of AI behind the ‘zombie internet’	The Guardian	AI Slop
3	Why doesn’t Facebook just ban AI slop like Shrimp Jesus?	Business Insider	AI Slop
4	Why the internet is filling up with nonsense ‘AI slop’	The Telegraph	AI Slop
5	Drowning in Slop	Intelligencer	AI Slop
6	How Long Will A.I.’s ‘Slop’ Era Last?	The New York Times	AI Slop
7	How Businesses Can Avoid AI Slop	Entrepreneur	Business
8	First Came ‘Spam.’ Now, With A.I., We’ve Got ‘Slop’	The New York Times	AI Slop
9	How ‘AI slop’ is creating a ‘zombie internet’	The Washington Post	AI Slop
10	Slop is the new name for unwanted AI-generated content	Simon Willison’s Weblog	AI Slop
11	Where Facebook’s AI Slop Comes From	404 Media	AI Slop
12	Side job, self-employed, high paid: behind the AI slop flooding TikTok and Facebook	The Conversation	AI Slop
13	Facebook’s Twisted Incentives Created Its AI Slop Era	Gizmodo	AI Slop
14	The MAGA Aesthetic Is AI Slop	The Atlantic	Politics
15	Hurricane Helene and the ‘Fuck It’ Era of AI-Generated Slop	404 Media	Politics
16	I’m Running Out of Ways to Explain How Bad This Is	The Atlantic	Culture
17	Coca-Cola’s New AI-Generated Holiday Ad Slammed as ‘Soulless’ and ‘Embarrassing’: ‘This Is Such Slop’	IGN	Advertising
18	Paramount skips voice actors, opting for shockingly bad AI slop in Novocaine film promo	TweakTown	Advertising
19	AI slop: what Labour, Spotify and Coca-Cola can teach us in 2025	Raconteur	Advertising

No.	Name	Source	Topic
20	A24 under fire for using AI-generated images to promote Civil War	MobileSyrup	Advertising
21	Activision “AI slop” Guitar Hero advert is market research for a fake game that doesn’t exist and maybe never will	Eurogamer	Advertising
22	Why Brands Are Embracing ‘Italian Brainrot’ to Go Viral on TikTok and Win Over Gen Z	NEWS10 ABC	Advertising
23	Willy Wonka event leaves bitter taste thanks to AI adverts	TheRegister	Fake Events
24	Chaos in Dublin as thousands turn up for AI ‘hoax’ Halloween parade that didn’t exist	The Independent	Fake Events
25	All Trick, No Treat: Dublin Crowds Turn Up for Halloween Parade That Wasn’t	The New York Times	Fake Events
26	Dublin: Halloween parade listing ‘mistake’ says website owner	BBC News	Fake Events
27	The Guy Behind the Fake AI Halloween Parade Listing Says You’ve Got It All Wrong	Wired	Fake Events
28	AI-Generated Slop Is Already In Your Public Library	404 Media	Publishing
29	Call of Duty Fans Give Black Ops 6’s Zombie Santa Loading Screen the Finger Amid ‘AI Slop’ Backlash	IGN	Games
30	Activision Reportedly Sold an AI-Generated Call of Duty: Modern Warfare 3 Cosmetic	IGN	Games
31	Call Of Duty Discloses AI Slop After Months Of Players Complaining	Kotaku	Games
32	PlayStation Exclusive Foamstars Has AI-Generated Art, Square Enix Confirms	IGN	Games
33	Square Enix President States the Company Will Be ‘Aggressive in Applying AI’	IGN	Games
34	Angry Birds Block Quest: Rovio’s latest title in the franchise has soft-launched for Android in the US	GamingonPhone	Games

B Appendix: Slop-related incidents

Disclaimer: this table was entirely generated using AI.

Incident	Category	Description	Example
Dublin “Hoax” Halloween Parade	Public Hoax / Fake Event	Thousands gathered in Dublin for a non-existent Halloween parade that was promoted by an AI-assisted article on a website. Irish police had to disperse the crowd.	AI-generated article on MySpiritHalloween.com.
Willy’s Chocolate Experience	Public Hoax / Fake Event	An event in Glasgow used fantastical AI-generated images for promotion but delivered a desolate, sparsely decorated warehouse. Children were left crying and police were called.	AI-generated promotional images and an “AI-generated gibberish” script.
Coca-Cola’s AI Holiday Ad	Advertising & Marketing	A fully AI-generated holiday ad faced widespread public backlash for being “soulless,” “creepy,” and replacing the work of human artists.	AI-generated holiday commercial featuring uncanny animals and trucks.
Spotify Wrapped 2024	Advertising & Marketing	The annual music summary was criticized as “AI-generated slop” due to bizarre, made-up genres and an impersonal AI podcast feature that users found inauthentic.	Made-up music genres (e.g., “Pink Pilates Princess Strut Pop”).
Paramount’s <i>Novocaine</i> Promotion	Advertising & Marketing	An Instagram promo for the film used a robotic AI voice for narration, leading to criticism for its low-effort, “shockingly bad” quality.	AI-generated voiceover and script for a movie promo reel.
A24’s <i>Civil War</i> Posters	Advertising & Marketing	The film company was criticized for using AI to create inaccurate promotional posters that did not depict actual scenes from the movie.	AI-generated movie posters, including one of soldiers firing on a giant swan.
Activision’s Fake Game Adverts	Advertising & Marketing	Ads for non-existent mobile games (<i>Guitar Hero Mobile</i> , etc.) used poor-quality AI art as part of a market research campaign to gauge interest.	AI-generated ads with flawed character art (e.g., a character missing half a face).


Incident	Category	Description	Example
“Italian Brainrot” Trend	Advertising & Marketing	Brands like Ryanair and Samsung adopted a surreal, low-effort AI meme trend to create chaotic, viral content intended to appeal to Gen Z.	AI-generated videos featuring nonsensical characters like “Tralalero Tralala.”
<i>Call of Duty: Black Ops 6</i> AI Assets	Video Games	Players identified AI-generated assets, most notably a loading screen of a zombified Santa Claus (“Necroclaus”) with six fingers. Activision later disclosed its use of AI.	“Necroclaus” loading screen with six fingers.
<i>Call of Duty: Modern Warfare 3</i> AI Cosmetic	Video Games	A report alleged that Activision sold a paid cosmetic bundle (“Yokai’s Wrath”) that was created with AI without disclosing it to customers.	The “Yokai’s Wrath” in-game cosmetic bundle.
<i>Foamstars</i> AI Album Art	Video Games	Square Enix confirmed using the AI tool Midjourney to create in-game album covers for its soundtrack, calling it a minor “experiment.”	In-game album art for the soundtrack.
<i>Angry Birds Block Quest</i> AI Art	Video Games	The soft-launched mobile game was criticized by fans for using generative AI for its cover art and other in-game assets.	Game cover art and various in-game assets.
Hurricane Helene Fake Photo	Political Dis-information	An AI image of a girl crying in a flood went viral after being shared by political figures to criticize the disaster response, even after they acknowledged it was fake.	Viral AI-generated image of a young girl holding a puppy in a flood.
The MAGA Aesthetic & AI Slop	Political Dis-information	Far-right influencers and Donald Trump’s accounts flooded social media with AI-generated images to mock opponents and create flattering, fantastical depictions of Trump.	AI images of Trump as a cowboy or a Steelers lineman.
Fake Taylor Swift Endorsements of Trump	Political Dis-information	A murky non-profit created and spread AI-generated images of Taylor Swift and her fans supporting Donald Trump, which he then shared on Truth Social.	AI-generated images of “Swifties for Trump.”

Incident	Category	Description	Example
Facebook’s “Shrimp Jesus” & Slop Farms	Low-Quality Content / Content Farms	Facebook was flooded with bizarre, viral AI images created by global “slop farms” to earn money from the platform’s creator bonus program.	Viral AI images of “Shrimp Jesus,” maimed veterans, and other engagement bait.
AI-Generated Books	Low-Quality Content / Content Farms	Low-quality, often dangerously inaccurate AI-generated books with fake authors appeared on Amazon and in public library digital catalogs.	Inaccurate mushroom foraging guides and medical advice books.
AI-Generated News & Articles	Low-Quality Content / Content Farms	Websites published AI articles with false and absurd information, such as a Microsoft Travel article listing a food bank as a tourist attraction.	A Microsoft Travel article about the “Ottawa food bank.”
<i>Late Night with the Devil</i> AI Art	Film & Television	The horror film was criticized for using low-quality AI images for its title cards, which featured anatomical inaccuracies.	Interstitial title card with a skeleton having poorly rendered fingers.
Amazon Prime Video AI Content	Film & Television	The streaming service used inaccurate AI-generated posters and thumbnails for classic films, as well as nonsensical, incomplete plot summaries.	Poster for <i>Nosferatu</i> ; thumbnail for <i>12 Angry Men</i> ; plot summary for <i>Dog Day Afternoon</i> .

C Appendix: BERTopic dendrogram

To improve readability, the plot is shown on the entirety of the next page.

Hierarchical Clustering



D Appendix: LLM-based thematic modelling

Completion markup was edited for readability; the text remains exactly the same.

Prompt:

use the qualitative method of thematic analysis on the following articles to answer the question ‘what is ai slop’ ### {{concatenated text of all documents}}

Completion:

Excellent. This is a perfect use case for thematic analysis. By systematically identifying, analyzing, and reporting patterns (themes) across this set of articles, we can construct a comprehensive, evidence-based answer to the question, “What is AI slop?”

Methodology

A qualitative thematic analysis was conducted on the provided articles. The process involved:

1. **Familiarization:** Reading and re-reading the articles to gain an in-depth understanding of the content.
2. **Coding:** Systematically identifying and labeling key concepts, ideas, and phenomena related to AI slop.
3. **Theme Development:** Collating the codes into broader, overarching themes that capture the essence of the data.
4. **Review and Refinement:** Reviewing the themes against the original data to ensure they are coherent, distinct, and well-supported.
5. **Synthesis:** Writing up the analysis, using direct quotes and examples from the articles to illustrate each theme.

Five primary themes emerged from the analysis, which together define the concept of “AI slop.”

—

Thematic Analysis: What is AI Slop?

Based on the provided articles, “AI slop” is a multifaceted phenomenon characterized by its poor quality, its economic and political motivations, its deceptive and harmful real-world impacts, and its contribution to a polluted and increasingly distrusted information ecosystem.

Theme 1: Defining Traits - Low Quality, Incoherent, and Uncanny

The most immediate characteristic of AI slop is its poor quality. It is consistently described as being carelessly produced, nonsensical,

and visually or textually flawed. This is not sophisticated deception, but rather mass-produced digital clutter.

* **Visual Flaws:** A common telltale sign is the “uncanny valley” effect, especially in AI-generated art. The articles point to a zombie Santa with “six fingers” in **Call of Duty**, AI-generated art for the movie **Nosferatu** that “does not resemble his look in the film,” and an AI image of a squirrel-like character “missing half a face” in a fake **Crash Bandicoot** ad.

* **Textual Gibberish:** The text associated with slop is often nonsensical or stilted. The Willy Wonka event featured an “AI-generated gibberish” script, and one AI-generated book review warned against a guide that contained “dangerous advice for anyone hoping to discern a lethal fungus from an edible one.” This extends to unintelligible captions like “catgacating” and “exarserdray lollipops” from the Willy Wonka promotional material.

* **Low-Effort & Soulless Nature:** The content is perceived as lacking human creativity and emotion. Coca-Cola’s AI-generated holiday ad was slammed as “‘soulless’ and ‘embarrassing’: ‘This is such slop.’” The result is content that is “cheap,” “shoddy,” and “devoid of any actual creativity.”

Theme 2: Motivations for Creation - Profit, Speed, and Scale

Slop is not created in a vacuum; it is the product of specific economic and efficiency-driven incentives. The primary motivation is rarely artistic or informational, but rather financial gain, speed, and the ability to produce content at an unprecedented scale.

* **Monetization & Engagement Farming:** The articles repeatedly link slop to revenue. Facebook’s “Creator Bonus Program” incentivizes creators in countries like India and Vietnam to churn out “bizarre AI spam” because “Facebook now pays you \$100 for 1,000 likes.” The owner of the fake Halloween parade website admitted his business model is “Google Ads and affiliate marketing,” and it’s “easy for us to generate revenue” from high-traffic holiday events.

* **Efficiency and Cost-Cutting:** Corporations use AI to drastically reduce production time and costs, even at the expense of quality. An AI studio founder touted that creating the Coca-Cola ad was “five times” faster than traditional methods. Activision allegedly laid off 2D artists and “forced” remaining concept artists to use AI. The use of AI is seen as a way to enhance “development productivity and achieve greater sophistication in our marketing efforts,” as stated by the president of Square Enix.

* **Low Barrier to Entry:** Generative AI has made it “trivial to automatically generate vast quantities of text or images.” This has

fueled a global gray-market economy where influencers on YouTube and Telegram teach others how to create slop as a “side job, self-employed, high-paid” venture.

Theme 3: Impact and Consequences - Misinformation, Deception, and Real-World Harm

Beyond being just digital junk, AI slop has tangible and often negative consequences. It deceives people, spreads misinformation, and can lead to real-world chaos and danger.

* **Hoaxes and Fake Events:** The most prominent example is the Dublin Halloween parade, where an “AI ‘hoax’” listing on a website led “thousands of people” to gather for an event that “never even existed.” Similarly, the “Willy’s Chocolate Experience” used AI-generated promotional images to sell tickets to a “desolate warehouse” event, leaving children crying and parents calling the police.

* **Dangerous Information:** The proliferation of AI-generated books has led to “dangerous advice” being sold on Amazon, such as mushroom foraging guides that could lead to poisoning. Public libraries are now inadvertently spending “taxpayer dollars to provide patrons with information they don’t realize is AI-generated.”

* **Political Weaponization:** Slop is used as a political tool to spread propaganda and sow distrust. During Hurricane Helene, an AI-generated image of a girl and a puppy was circulated by political figures to suggest government failure. One pundit justified it, saying, “I’m leaving it because it is emblematic of the trauma and pain people are living through right now,” typifying what one article calls the “‘Fuck It’ Era of AI slop,” where emotional appeal trumps truth.

Theme 4: The ‘Slop’ Framework - A Name for a Polluted, ‘Post-Truth’ Internet

The emergence of the term “slop” itself is a key theme. It provides a lexicon to describe a new kind of internet pollution and the broader societal shift it represents.

* **A New Term for a New Problem:** The articles frame “slop” as the successor to “spam.” As one developer puts it, “the way that ‘spam’ became the term for unwanted emails, ‘slop’ is going in the dictionary as the term for unwanted AI generated content.” This naming is important because “it gives people a concise way to talk about the problem.”

* **The Zombie Internet:** Slop contributes to what is called the “zombie internet,” a space where “a mix of bots, humans and accounts that were once humans but aren’t any more mix together to form a disastrous website where there is little social connection at all.” This cheap, automated content clutters feeds and search results, “making it harder for them to find accurate information.”

* **Erosion of Reality and Trust:** The ultimate consequence of slop is a breakdown in shared reality. When politicians and influencers knowingly share fake images because they are “emblematic” of a feeling, it signals a move to a “post-reality” where truth is secondary. This erodes trust not just in online content, but in institutions like news media, government agencies (FEMA), and science. As one article concludes, “Americans are divided not just by political beliefs but by whether they believe in a shared reality—or desire one at all.”

Conclusion

In summary, the thematic analysis of these articles reveals that AI slop is low-quality, often incoherent, and uncanny content mass-produced by generative AI. It is primarily motivated by the pursuit of profit, speed, and scale, both by individual “sloppers” and large corporations. Its impact is far from benign, leading to real-world deception, dangerous misinformation, and political manipulation. Finally, the term “slop” itself is a new framework for understanding the pollution of the internet, a “zombie internet” where the lines between human and machine, and truth and fiction, are increasingly and deliberately blurred.“

Prompt (another chat):

{{previous response}} ### summarize the themes into a bullet point list with a brief description for each

Completion:

Of course. Here is a summary of the four primary themes in a bullet point list with a brief description for each.

Summary of Themes Defining “AI Slop”

*** Theme 1: Defining Traits - Low Quality and Uncanny**

This theme covers the fundamental characteristics of the content itself. Slop is defined by its poor quality, visual flaws (like extra fingers or distorted faces), nonsensical text, and a “soulless” or cheap feel that betrays its non-human origin.

*** Theme 2: Motivations - Profit, Speed, and Scale**

This theme explains the reasons behind slop’s creation. It is driven by economic incentives, such as individuals monetizing engagement through ads (“engagement farming”), and corporate desires to cut costs and dramatically accelerate content production by replacing human labor.

*** Theme 3: Impact and Consequences - Deception and Real-World Harm**

This theme highlights that slop is not harmless. It has tangible negative effects, from deceiving the public with fake events and promotions (the Willy Wonka experience) to spreading dangerous misinformation (faulty foraging guides) and being weaponized for political propaganda.

*** Theme 4: The 'Slop' Framework - A Polluted Information Ecosystem**

This theme addresses the broader context and consequences. The term “slop” itself provides a new name for this form of digital pollution, which contributes to a “zombie internet” cluttered with automated content. Ultimately, it erodes public trust and blurs the lines between truth and fiction in a “post-truth” environment.

E Appendix: Sources of slop

Source	Website	Operational
NewsGuard	famadillo.com	X
NewsGuard	wavefunction.info	
NewsGuard	scoopearth.com	X
NewsGuard	filthylucre.com	X
NewsGuard	dailybusinesspost.com	X
NewsGuard	getintoknowledge.com	X
NewsGuard	bestbudgetusa.com	
NewsGuard	harmonyhustle.com	
NewsGuard	historyfact.in	
NewsGuard	countylocalnews.com	X
NewsGuard	tnewsnetwork.com	
NewsGuard	celebritiesdeaths.com	
NewsGuard	Biz Breaking News	
NewsGuard	News Live 79	
NewsGuard	Market News Reports	
DoubleVerify	espn24.co.uk	
DoubleVerify	nbc sport.co.uk	X
DoubleVerify	cbsnewz.com	
DoubleVerify	cbsnews2.com	X
DoubleVerify	bbc sportss.co.uk	
DoubleVerify	foxnigeria.com.ng	X
DoubleVerify	bbc sportz.com	
DoubleVerify	247bbcnews.com	
WIRED	thehairpin.com	X
WIRED	antoniocarluccio.com	
WIRED	pope2you.net	X
WIRED	trumpplaza.com	X
WIRED	thefrisky.com	X